

Highlights

Pseudo-3D Heightfield Reconstruction from Endoscopic Video via Normal-Guided Multi-View Optimization and Metric Scale Gauge Injection

Ranyang Li, Wufeng Liu, Xiaodong Guo, Nan Wei, Chao Fan

- Heightfield normal field resolves two-fold planar homography ambiguity.
- Multi-stride BA adds links across non-adjacent keyframes.
- Board gauge sets millimetre scale when the pattern is visible.
- Board-gauge median ATE is 4.49 mm on 13 real sequences.
- Calibrated stereo variant achieves 0.0285 AbsRel on public EndoNeRF.

Pseudo-3D Heightfield Reconstruction from Endoscopic Video via Normal-Guided Multi-View Optimization and Metric Scale Gauge Injection

Ranyang Li^{a,*}, Wufeng Liu^a, Xiaodong Guo^a, Nan Wei^b, Chao Fan^a

^a*School of Artificial Intelligence and Big Data, Henan University of Technology, Zhengzhou 450001, China.*

^b*Henan Provincial People's Hospital, Zhengzhou University, Zhengzhou 450003, China.*

Abstract

Recovering visualizable, metrically scaled 3D surfaces from monocular endoscopic video is a critical prerequisite for intra-operative surgical navigation and anatomical measurement. Existing methods struggle with three fundamental dilemmas: feed-forward foundation models suffer from severe metric scale distortion ($8\times$ – $647\times$) and planar collapse; classical monocular visual geometry suffers from compounding dead-reckoning drift; and small-baseline planar homography decomposition is corrupted by a mathematical two-fold branch ambiguity. To resolve these challenges, we propose a unified geometric framework integrating pseudo-3D heightfield surface representation, normal-guided branch disambiguation, multi-stride bundle adjustment, and metric scale gauge injection. First, a bilateral gradient-anchored 2.5-D heightfield mesh provides a well-conditioned continuous surface and an operative physical normal field $\mathbf{n}_{hf}$. Second, we introduce a heightfield normal veto that selects the homography branch whose plane normal agrees with $\mathbf{n}_{hf}$. On the evaluated pairs this choice has median directional agreement 0.984 with the reference translation. Third, a multi-stride co-visibility graph across non-adjacent keyframes feeds a joint Lie-algebra bundle adjustment on $SE(3)$, forming closed geometric loops that suppress sequential drift. Fourth, a metric scale gauge restores true millimetre dimensions from physical calibration landmarks. On the 13 real sequences in which a 3 mm chessboard is visible, full-gauge injection reaches a

*Corresponding author

Email address: lry@haut.edu.cn (Ranyang Li)

median ATE of 4.49 mm, against 11.88 mm for VGGT-1B and 41.98 mm for ORB-SfM on those same sequences. Without that gauge, multi-view BA has a 17-sequence median ATE of 29.88 mm: better than ORB-SfM (40.29 mm) and not better than VGGT-1B (16.86 mm). On public EndoNeRF, the calibrated stereo variant reaches an overall AbsRel of 0.0285 (FoundationStereo 0.0321). The result is an accurate, metrically anchored 3D representation robust to endoscopic geometry.

Keywords: Endoscopy, Reconstruction, Visualization, Laparoscopy, Stereoscopy, Heightfield

1. Introduction

Recovering visualizable, metrically scaled 3D surfaces from monocular endoscopic video is a core imaging problem for intra-operative navigation, augmented-reality overlay, and quantitative surgical assessment in minimally invasive surgery (MIS) [1, 2].

5 Unlike CT or MRI, the endoscope is already in the sterile field and supplies the only continuous intra-operative view; converting that video into a metric surface therefore determines whether graphical overlays and image-guided measurements can be trusted. A missing or wrong metric scale has direct consequences: estimating the safe instrument-to-tissue standoff that avoids serosal injury, measuring lesion diameter for

10 resection-margin planning, and quantifying tissue deformation during suturing. These tasks require *metric* (not merely relative) depth, yet most existing pipelines recover only up-to-scale geometry, and the monocular setting removes the baseline cue that stereo reconstruction relies on. General-purpose 3D vision has advanced rapidly, yet monocular endoscopic reconstruction remains plagued by four fundamental dilemmas.

15 First, smooth tissue surfaces offer repetitive texture and few discriminative features. Second, specular reflections from wet tissue violate the photometric consistency that dense reconstruction assumes. Third, small-baseline planar motion induces an analytical two-fold branch ambiguity in homography decomposition, causing conventional front-ends to reverse translation directions. Fourth, sequential dead-reckoning in

20 SLAM/SfM causes severe compounding drift, while feed-forward natural-image foundation models (e.g., DUS3R [3], MAST3R [4], VGGT [5]) collapse into flat slabs or

cones ($\lambda_3/\lambda_1 \leq 0.22$) with arbitrary scale wandering ($8\times\text{--}647\times$).

Classical structure-from-motion (SfM) methods [6, 7] fail on such data: on our benchmark COLMAP registers only a fraction of frames in textureless regions and
25 produces cone-shaped clouds ($\lambda_3/\lambda_1 < 0.13$), and dense SLAM [8, 9] suffers specular divergence. Feed-forward networks that excel on natural images degrade on endoscopic footage because networks trained on internet datasets [10] encode matching and scale priors that endoscopic illumination and geometry violate.

The thesis of this paper is that robust, metrically faithful reconstruction requires
30 combining a *well-conditioned physical surface embedding* with *normal-guided branch disambiguation*, *multi-stride bundle adjustment*, and *metric scale gauge injection*. We construct the surface embedding as a bilateral gradient-anchored 2.5-D heightfield mesh, avoiding learned natural-image biases. Crucially, the heightfield geometry provides not only a visualizable rendering primitive, but also an explicit physical surface normal
35 field $\mathbf{n}_{hf}$. When decomposing planar homographies between frames, this normal field selects the branch aligned with the surface and, on the evaluated pairs, removes the 180° translation flip. Furthermore, instead of relying on fragile Markovian dead-reckoning, we construct a multi-stride co-visibility graph across non-adjacent keyframes and perform joint Bundle Adjustment (BA) over $\text{SE}(3)$ Lie-algebra camera poses and
40 triangulated 3D points. Finally, physical calibration landmarks provide an exact metric gauge injection, converting dimensionless relative translations into true millimetre dimensions.

When a calibrated stereo pair is available, a learned disparity network replaces the gradient proxy with metric depth, and the heightfield is retained as a shape regularizer
45 rather than discarded.

The core contributions of this work are summarized as follows:

- **Physical surface embedding with normal prior.** We formulate a bilateral gradient-anchored 2.5-D heightfield that simultaneously serves as the display surface, the pose registration object, and the dense matching substrate. The mesh generates a
50 continuous operative normal field $\mathbf{n}_{hf}$ that prevents geometric cone collapse ($\lambda_3/\lambda_1 = 0.42$).

Table 1: Capability comparison of 3D reconstruction methods on endoscopic data. $\checkmark$: fully supported; Δ : partially supported or limited; $\times$: not supported. Ours supports both lightweight CPU-only deployment (Classical mode) and GPU-accelerated multi-view bundle adjustment (Ours-BA mode).

Capability	ORB-SfM	COLMAP	VGGT	DUSt3R	Ours (Classic)	Ours-BA
Dense recon.	$\times$	$\times$	$\checkmark$	$\checkmark$	Δ	$\checkmark$
Metric scale	Δ	$\times$	$\times$	$\times$	$\checkmark$	$\checkmark$
Normal 3D shape	$\times$	$\times$	$\times$	$\times$	Δ	$\checkmark$
Specular robust.	$\times$	$\times$	Δ	Δ	Δ	$\checkmark$
CPU-only deploy.	$\checkmark$	$\checkmark$	$\times$	$\times$	$\checkmark$	$\times$
Repetitive texture	$\times$	$\times$	$\times$	$\times$	Δ	$\checkmark$

- **Normal-guided homography branch disambiguation.** We analyze the analytical two-fold physical ambiguity in planar homography decomposition under small endoscopic baselines, and propose a heightfield normal veto criterion that selects the true camera motion branch without external sensors.
- **Multi-stride co-visibility graph and joint bundle adjustment.** We replace drifting sequential pose chaining with a multi-stride co-visibility graph across multi-frame skips ($s \in \{1, 2, 3, 4, 5\}$). A joint Lie-algebra Bundle Adjustment on $SE(3)$ with Huber loss resolves loop closures and suppresses long-range drift.
- **Metric scale gauge injection and a 17-sequence benchmark.** When a 3 mm chessboard is visible, gauge injection anchors translation in millimetres. On the 13 real sequences the full-gauge median ATE is 4.49 mm (VGGT-1B 11.88 mm and ORB-SfM 41.98 mm on the same 13). Marker-free multi-view BA, scored on all 17 sequences, has median ATE 29.88 mm, which beats ORB-SfM (40.29 mm) and does not beat VGGT-1B (16.86 mm).

Table 1 compares the proposed method with representative approaches on dense reconstruction, metric scale, 3D shape, and CPU-only operation.

2. Related Work

Our research is related to endoscopic 3D reconstruction and visualization, classical
70 and learning-based structure-from-motion, and calibration-pattern-assisted methods in
medical imaging. We briefly review representative works in each of these categories.

Endoscopic 3D reconstruction. Early work on endoscopic 3D reconstruction
focused on stereo endoscopes [1, 2], which provide direct depth from binocular dis-
parity, but most surgical endoscopes are monocular. Shape-from-shading (SfS) [11]
75 recovers surface orientation under Lambertian assumptions but fails under the non-
Lambertian reflections prevalent in wet tissue. Structure-from-motion approaches [12]
rely on feature correspondences but yield sparse reconstructions in texture-poor images.
Learning-based monocular depth estimation [13, 14] generalizes poorly to endoscopic
scenes; recent surgical methods [15] incorporate temporal consistency and surgical
80 priors yet still produce only relative depth. Classical stereo matching, in particular
semi-global matching (SGM) [16], recovers dense disparity from rectified stereo pairs
and remains a strong, training-free baseline for endoscopic depth when a stereo en-
doscope is available; we therefore adopt it as both a depth source and a comparison
baseline throughout this paper. Reconstructed surfaces are useful only if they can be
85 rendered and registered: prior work on intra-operative augmented reality and image-
guided intervention [17, 1] therefore treats visualization quality—well-formed rather
than collapsed geometry—as a requirement, not a post-processing detail.

Classical structure-from-motion. ORB-SLAM [6] uses ORB features [18] for
real-time tracking; COLMAP [7] employs SIFT [19] with incremental reconstruction
90 and bundle adjustment; LSD-SLAM [20] performs direct semi-dense reconstruction.
All three rely on discriminative features or photometric consistency, both scarce in
endoscopic scenes. Dense SLAM (DTAM [8], ElasticFusion [9]) similarly requires
sufficient texture. RGB-D methods (KinectFusion [21], BundleFusion [22]) are inap-
plicable to monocular endoscopy.

95 **Feed-forward 3D reconstruction networks.** Recent foundation models for 3D
reconstruction have shifted from classical optimization to feed-forward regression.
DUST3R [3] predicts dense paired pointmaps with global alignment; MAST3R [4] adds

local feature matching and metric heads; MUST3R [23] generalizes this to multi-view video via memory-buffered attention; VGGT [5] (CVPR 2025 Best Paper) and its successor VGGT- Ω predict camera poses, depth maps, and 3D point clouds in a single pass using a Vision Geometry Grounded Transformer; and Reloc3R [24] optimizes camera relocalization from dense point correspondences. Despite their success on internet-scale natural scenes, these feed-forward architectures exhibit severe failure modes when applied directly to monocular endoscopic video: (1) *Scale wandering*: because monocular feeds lack physical baselines, their predicted translations wander arbitrarily (e.g., $8\times$ – $647\times$ scale factors on endoscopic phantoms); (2) *Geometric collapse*: under repetitive surgical texture and specular wet surfaces, self-attention maps produce planar slabs or narrow cones ($\lambda_3/\lambda_1 < 0.20$), failing to preserve 3D anatomical volume; (3) *Out-of-distribution drift*: as evidenced in our 120-epoch training experiments, feed-forward regressors struggle to extrapolate to unseen camera distances and large sweep motions without explicit geometric constraints.

Calibration-pattern-based methods and metric gauge in MIS. Calibration patterns (e.g., planar chessboards) have long been used for camera calibration [25, 26], instrument pose tracking [27], and augmented reality overlays [17]. However, classical methods treat pattern corners merely as sparse feature points for PnP tracking, yielding no dense anatomical surface. In our framework, we re-purpose physical patterns not as the reconstructed target, but as a *metric scale gauge* that injects exact physical scale into multi-stride bundle adjustment, while the dense surgical anatomy is reconstructed via the gradient-anchored heightfield.

Summary of gaps. Taken together, monocular endoscopic 3D reconstruction requires simultaneously solving geometric collapse, small-baseline planar ambiguity, sequential drift, and metric scale loss. Classical SfM yields sparse or disconnected trajectories; feed-forward networks collapse geometrically without metric scale; and conventional homography decomposition produces branch flips. The proposed method bridges these gaps by coupling a physical heightfield surface with normal-guided branch disambiguation, multi-stride bundle adjustment, and metric scale gauge injection.

3. Method

3.1. Problem Formulation

Given a monocular endoscopic video sequence $\{I_t\}_{t=0}^{N-1}$ of N frames, our goal is to
130 recover a dense, metrically scaled 3D point cloud $\mathcal{P} = \{(\mathbf{X}_k, \mathbf{c}_k)\}_{k=1}^K$ of the surgical
scene, where $\mathbf{X}_k \in \mathbb{R}^3$ denotes a 3D point and $\mathbf{c}_k \in [0, 255]^3$ its associated color.
The reconstructed cloud, together with the per-frame heightfield mesh and an optional
confidence map, is intended as a visualizable geometric layer for intra-operative graphics
and image-guided measurement. When a planar calibration pattern (e.g., a chessboard
135 marker) with known geometry is present, it provides optional geometric anchors that
constrain the pipeline, as detailed in the following sections.

Notation.. We denote camera poses as $\mathbf{T}_t \in \text{SE}(3)$ (camera-to-world), the camera
intrinsic matrix as $\mathbf{K} \in \mathbb{R}^{3 \times 3}$, and the chessboard corner set as $\{\mathbf{u}_{t,k}\}_{k=1}^M$ for frame t .
The coordinate system uses (u, v) for pixel coordinates and h for the pseudo-height
140 derived from image gradients. A complete reference of mathematical symbols and
dimensional definitions is summarized in Supplementary Material, Table S1.

3.2. Pipeline Overview

The proposed framework processes monocular endoscopic video through four se-
quential stages (Fig. 1). First (Section 3.3), each frame is converted into a 2.5-D
145 heightfield mesh from truncated image gradients and a region-of-interest mask derived
from a calibration pattern or other detectable landmarks when available; the mesh cap-
tures the topology of the visible tissue surface and can be rendered directly. Second
(Section 3.4), inter-frame poses are estimated by registering consecutive heightfields
with FPFH coarse alignment (PCA or feature-matched RANSAC) and point-to-point
150 ICP, gated by an online quality measure. Third (Section 3.7), the estimated poses
establish dense correspondences through epipolar-constrained geometric matching on
the heightfield. Fourth (Section 3.7), per-frame reconstructions are triangulated and
fused into a globally consistent metric point cloud. When a calibrated stereo pair is
available, an optional end-to-end variant (Section 3.8) recovers metric depth through




Figure 1: Overview of the proposed pseudo-3D endoscopic reconstruction framework with normal-guided multi-view optimization and metric scale gauge injection. **Stage 1 (Heightfield and normal field construction)**: bilateral gradient filtering, percentile truncation, and ROI masking generate a continuous 2.5-D heightfield mesh $\mathcal{H}_t$ and dominant physical normal field $\mathbf{n}_{hf}$. **Stage 2 (Normal-guided homography disambiguation)**: relative planar homography $\mathbf{H}$ between frames is decomposed, and the true physical camera motion branch $\mathbf{T}_{ij}$ is selected via the heightfield normal veto rule $k^* = \arg \max |\mathbf{n}_k^T \bar{\mathbf{n}}_{hf}|$, eliminating 180° translation flips. **Stage 3 (Multi-stride graph and bundle adjustment)**: multi-stride co-visibility edges ($s \in \{1..5\}$) form closed geometric loops on $\text{SE}(3)$, optimized jointly with 3D points under robust Huber loss to eliminate dead-reckoning drift. **Stage 4 (Metric gauge injection and dense fusion)**: physical standoff distance d_t injects 1:1 metric scale (s_{sq}), followed by multi-view consistency filtering and dense point cloud fusion into $\mathcal{P}$. **Legend**: numbered circles denote stage index; light boxes are inputs, filled boxes are processing steps, dark boxes are outputs; solid arrows denote primary data flow; dashed arrows denote auxiliary data flow; equation annotations summarise the core operation of each stage.

155 a learned disparity network regularized toward geometrically consistent 3D structure.
The classical monocular pipeline remains the primary method.

3.3. Pseudo-3D Heightfield Mesh Construction

Image gradients encode the topological structure of the surgical scene: strong gradients mark structural edges, while weak gradients from homogeneous regions are
160 suppressed by percentile truncation.

Gradient computation and truncation.. For each frame I_t of spatial resolution $H \times W$ (rows $\times$ columns), we compute the gradient magnitude map $g_t(u, v) = \|\nabla I_t(u, v)\| = \sqrt{(\partial I_t / \partial u)^2 + (\partial I_t / \partial v)^2}$. To suppress noise and spurious gradients from specular reflections, we apply percentile-based truncation: $\hat{g}_t(u, v) = g_t(u, v)$ if $g_t(u, v) \geq \tau_g$,
165 and 0 otherwise, where $\tau_g = \text{Perc}_{\tau_{\%}}(g_t)$ denotes the $\tau_{\%}$ -th percentile of the gradient distribution. With $\tau_{\%} = 68$, the upper portion of the gradient distribution is retained. The percentile is a fixed operating point, not a contribution of this paper.

Region-of-interest mask. When detectable landmarks such as a calibration pattern are present, they provide a binary region-of-interest mask $M_t(u, v) \in \{0, 1\}$ (e.g.,
170 detected via standard corner detection [28] for chessboard patterns). The mask serves two purposes: (1) it suppresses background regions that would otherwise dominate the heightfield and corrupt ICP convergence; and (2) it defines the reliable region for subsequent pose quality evaluation. The truncated gradient map is masked as $\tilde{g}_t(u, v) = \hat{g}_t(u, v) \cdot M_t(u, v)$. This masking ensures that only gradient responses
175 within the region of interest contribute to the heightfield, removing the influence of irrelevant tissue and instrument regions. The interface is deliberately source-agnostic: M_t may come from a chessboard hull, a contour fallback, or a semantic tissue probability in $[0, 1]$. Instruments and specular blobs should be excluded from the tissue heightfield rather than reconstructed as a second depth layer, which a single 2.5-D mesh cannot
180 represent.

Heightfield mesh generation. The masked gradient map $\tilde{g}_t$ (normalised to $[0, 255]$) is lifted to a 2.5-D heightfield mesh $\mathcal{H}_t$ by constructing a regular grid with cell size c_{cell} pixels (default: 3). Each grid vertex (u, v) is assigned a height value $h(u, v) = z_s \cdot \tilde{g}_t(u, v)/255$, where $z_s = 0.15 \max(H, W)$ is the height scaling factor so that the
185 maximum height is comparable to image dimensions for numerical stability. Adjacent vertices are connected via a regular grid triangulation, yielding a manifold surface $\mathcal{H}_t = (\mathcal{V}_t, \mathcal{F}_t)$ with $|\mathcal{V}_t| \approx HW/c_{\text{cell}}^2$ vertices.

Physical surface normal field. In addition to height coordinates, each vertex (u, v) on the triangular mesh defines a continuous surface normal $\mathbf{n}(u, v)$ computed from
190 adjacent facet normals:

$$\mathbf{n}_{hf}(u, v) = \frac{\sum_{j \in \mathcal{F}(u, v)} \mathbf{n}_j}{\left\| \sum_{j \in \mathcal{F}(u, v)} \mathbf{n}_j \right\|_2}, \quad (1)$$

where $\mathcal{F}(u, v)$ denotes the set of triangular facets sharing vertex (u, v) , and $\mathbf{n}_j$ is the cross-product normal of facet j . In surgical cavities and anatomical phantom setups, the camera optical axis maintains a consistent viewing inclination relative to the underlying tissue bed. The dominant orientation $\bar{\mathbf{n}}_{hf} = \text{median}_{(u, v)} \mathbf{n}_{hf}(u, v)$ serves as an operative
195 physical normal prior that constrains inter-frame geometry.

3.4. Normal-Guided Planar Homography Disambiguation

When endoscopic cameras observe planar surgical phantoms, liver surfaces, or anatomical planes, points on a planar structure $\mathbf{n}^\top \mathbf{X} = d$ transform between views i and j via a homography matrix $\mathbf{H} \in \mathbb{R}^{3 \times 3}$:

$$\mathbf{p}_j \sim \mathbf{H} \mathbf{p}_i, \quad \mathbf{H} = \mathbf{K} \left(\mathbf{R}_{ij} + \frac{\mathbf{t}_{ij}}{d_i} \mathbf{n}_i^\top \right) \mathbf{K}^{-1}, \quad (2)$$

where $\mathbf{R}_{ij} \in \text{SO}(3)$ is the relative rotation, $\mathbf{t}_{ij} \in \mathbb{R}^3$ is the translation vector, $\mathbf{n}_i$ is the unit surface normal in the reference camera frame, and d_i is the orthogonal distance from camera i to the plane. The normalized homography $\tilde{\mathbf{H}} = \mathbf{K}^{-1} \mathbf{H} \mathbf{K}$ is estimated from robust feature correspondences (e.g., deep LoFTR matches or SIFT) via RANSAC.

The two-fold physical ambiguity. Analytical decomposition of $\tilde{\mathbf{H}}$ via singular value decomposition [29] yields four mathematical solutions, of which exactly two satisfy the positive cheirality constraint, denoted as physical candidates $\mathcal{S}_1 = (\mathbf{R}_1, \mathbf{t}_1/d, \mathbf{n}_1)$ and $\mathcal{S}_2 = (\mathbf{R}_2, \mathbf{t}_2/d, \mathbf{n}_2)$. Under narrow baselines typical of endoscopic videos ($\|\mathbf{t}\| \ll d$), both solutions yield almost identical reprojection errors on image features: $\|\mathbf{p}_j - \pi(\mathbf{H}_1 \mathbf{p}_i)\| \approx \|\mathbf{p}_j - \pi(\mathbf{H}_2 \mathbf{p}_i)\| < 1.0$ px. Standard heuristic rules—such as selecting the solution with the largest number of positive triangulated points or the minimal rotation—frequently fail (mis-selection rate $> 25\%$), choosing an unnatural fronto-parallel branch where the estimated translation vector is reversed or biased by 90° .

Heightfield normal veto. We resolve this ambiguity by evaluating the directional alignment between the candidate plane normals $\{\mathbf{n}_1, \mathbf{n}_2\}$ and the physical surface normal prior $\tilde{\mathbf{n}}_{hf}$ derived from the pseudo-3D heightfield mesh (Eq. (1)):

$$k^* = \arg \max_{k \in \{1,2\}} |\mathbf{n}_k^\top \tilde{\mathbf{n}}_{hf}|. \quad (3)$$

The selected solution $(\mathbf{R}_{k^*}, \mathbf{t}_{k^*}, \mathbf{n}_{k^*})$ is the decomposition whose plane normal is closer to $\tilde{\mathbf{n}}_{hf}$. This is a selection rule, not a proof that the selected branch is the true motion for every homography. On pairs with significant camera displacement, the selected candidate has median directional agreement $|\mathbf{t}_{k^*}^\top \mathbf{t}_{gt}| = 0.984$, and the opposite translation sign is not chosen.

3.5. Multi-Stride Co-Visibility Graph and Joint Bundle Adjustment

Conventional monocular visual odometry integrates relative poses in a sequential Markovian chain ($\mathbf{T}_{0 \rightarrow t} = \prod_{k=0}^{t-1} \mathbf{T}_{k \rightarrow k+1}$), which causes compounding dead-reckoning drift over extended sequences. To overcome this, we construct a *multi-stride co-visibility graph* and optimize camera poses and 3D scene points jointly via non-linear Bundle Adjustment.

Multi-stride graph construction.. Given N keyframes, we establish co-visibility edges across multi-frame temporal skips: $\mathcal{E} = \{(i, j) \mid 1 \leq j - i \leq S_{\max}, 0 \leq i < j < N\}$, where $S_{\max} \in \{3, 4, 5\}$ is the maximum co-visibility stride. For each edge $(i, j) \in \mathcal{E}$, deep correspondences $\mathcal{M}_{ij} = \{(\mathbf{u}_i^{(m)}, \mathbf{u}_j^{(m)})\}$ are extracted using LoFTR [30] and filtered via homography RANSAC. Multi-view tracks are formed by chaining pairwise matches across overlapping frames, and initial 3D point positions $\mathbf{X}_k \in \mathbb{R}^3$ are triangulated via multi-view linear SVD.

Lie-algebra bundle adjustment formulation.. We parameterize each camera pose $\mathbf{T}_i \in \text{SE}(3)$ (mapping world points to the camera frame) using $\mathfrak{se}(3)$ Lie-algebra twist coordinates $\xi_i = (\omega_i^\top, \mathbf{v}_i^\top)^\top \in \mathbb{R}^6$ via the exponential map:

$$\mathbf{T}_i = \exp(\xi_i^\wedge) = \begin{bmatrix} \exp(\omega_i^\wedge) & \mathbf{V}(\omega_i)\mathbf{v}_i \\ \mathbf{0}^\top & 1 \end{bmatrix}, \quad (4)$$

with the first camera fixed at the world origin ($\xi_0 = \mathbf{0}$). Here $\omega^\wedge$ is the 3×3 skew-symmetric matrix of ω , and $\mathbf{V}(\omega)$ is the left Jacobian of $\text{SO}(3)$,

$$\mathbf{V}(\omega) = \mathbf{I} + \frac{1 - \cos \theta}{\theta^2} \omega^\wedge + \frac{\theta - \sin \theta}{\theta^3} (\omega^\wedge)^2, \quad \theta = \|\omega\|, \quad (5)$$

with $\mathbf{V}(\mathbf{0}) = \mathbf{I}$. The joint objective over all free camera poses $\{\xi_i\}_{i=1}^{N-1}$ and 3D scene points $\{\mathbf{X}_k\}_{k=1}^K$ is formulated as:

$$\min_{\{\xi_i\}, \{\mathbf{X}_k\}} \sum_{(i,k) \in \mathcal{O}} \rho_{\text{Huber}} \left(\|\mathbf{u}_{i,k} - \pi(\mathbf{K} \exp(\xi_i^\wedge) \mathbf{X}_k)\|_2 \right) + \lambda_{\text{cheir}} \sum_{(i,k) \in \mathcal{O}} \max(0, d_{\min} - z_{i,k})^2, \quad (6)$$

where $\mathcal{O}$ is the set of all 2D-3D observation pairs, $\rho_{\text{Huber}}(\cdot)$ is the robust Huber penalty with threshold $\delta_H = 3.0 \text{ px}$, $z_{i,k} = [\exp(\xi_i^\wedge) \mathbf{X}_k]_z$ is the depth in camera frame i ,

and the second term enforces a positive cheirality safety barrier ($d_{\min} = 10.0$ mm). The objective is optimized on GPU via AdamW with cosine learning rate annealing, converging in 400–800 iterations. Because edges span across multi-frame skips, the resulting graph forms dense triangle and polygon closures that distribute local tracking errors and eliminate trajectory drift.

3.6. Metric Scale Gauge Injection

Monocular reconstruction from visual correspondences is fundamentally scale-ambiguous: the translation vector recovered from homography or essential matrix decomposition represents a dimensionless ratio $\hat{\mathbf{t}} = \mathbf{t}/d$. Feed-forward models (such as VGGT and DUSt3R) predict scale-less geometry that wanders by orders of magnitude across sequences ($8\times$ – $647\times$).

Physical scale recovery. When a physical reference of known pitch s_{sq} (e.g., standard 3.0 mm calibration squares or surgical fiducials) is visible in the endoscopic field of view, we inject an exact physical metric gauge. Let p_{sq} be the measured pixel pitch of adjacent pattern corners, and $f = (f_x + f_y)/2$ the focal length in pixels. The standoff is $d_t = f \cdot s_{\text{sq}}/p_{\text{sq}}$. If the decomposition returns an up-to-scale translation $\hat{\mathbf{t}}$, the metric translation is $\mathbf{t}_{\text{metric}} = d_t \hat{\mathbf{t}}$.

Sparse vs. full gauge propagation. In practice, full physical targets may not be visible in every surgical frame. Our framework supports two operating modes:

- **Sparse Gauge Mode:** The metric gauge is anchored on a sparse subset of keyframes (e.g., every 5th or 10th frame, or only when the pattern is fully unoccluded). The metric scale is then propagated to all intermediate frames through the multi-stride bundle adjustment graph (Section 3.5), holding scale drift to within $1.03\times$ – $1.10\times$.
- **Full Gauge Mode:** When the pattern is visible in the frames being scored, those frames receive metric scale constraints. On the 13 real sequences this median ATE is 4.49 mm. It is not sub-5 mm on every sequence: `traj_015507` remains 80.25 mm.

3.7. Dense Epipolar Matching and Multi-Frame Fusion

To reconstruct a *dense* point cloud beyond sparse landmarks or feature tracks, we establish pixel-level correspondences between frames using heightfield-based geometric feature matching with epipolar constraints.

275 *7D geometric feature extraction..* For each valid grid point on the heightfield, we construct a 7-dimensional feature vector:

$$\mathbf{f}(u, v) = \left[\underbrace{\frac{u}{W}, \frac{v}{H}}_{\text{position}}, \underbrace{\frac{h}{z_s}}_{\text{height}}, \underbrace{\frac{n_x}{\|\mathbf{n}\|}, \frac{n_y}{\|\mathbf{n}\|}, \frac{n_z}{\|\mathbf{n}\|}}_{\text{normal}}, \underbrace{\frac{c}{255}}_{\text{color}} \right]^\top, \quad (7)$$

combining normalized position $(u/W, v/H)$, height h/z_s , surface normal $(n_x, n_y, n_z)^\top / \|\mathbf{n}\|$, and colour $c/255$ (where $c \in [0, 255]$ is the grey-scale intensity at the vertex). The surface normal is estimated from the heightfield via finite differ-
280 ences, providing geometric context that disambiguates repetitive texture patterns.

RT-predicted search window.. Given the estimated relative transformation $\mathbf{T}_{t+1,t}$ from Section 3.4, which is defined in the heightfield space, we predict the location of each source grid point in the target frame via $[u', v', h', 1]^\top = \mathbf{T}_{t+1,t} [u, v, h, 1]^\top$, where (u, v, h) is the source vertex height and (u', v', h') its predicted counterpart. This
285 provides a tight search prior that reduces the matching search space from the full image to a local window of radius r (default: $r=25$ pixels) around the predicted location.

Epipolar constraint.. For frames with known camera intrinsics and estimated pose, we compute the fundamental matrix $\mathbf{F} = \mathbf{K}^{-\top} [\mathbf{t}]_\times \mathbf{R} \mathbf{K}^{-1}$, where $\mathbf{R}$ and $\mathbf{t}$ are the relative rotation and translation and $[\mathbf{t}]_\times$ is the skew-symmetric matrix of $\mathbf{t}$, and enforce
290 the epipolar band constraint $|\mathbf{u}'^\top \mathbf{F} \mathbf{u}| < \epsilon_{\text{epi}}$ (with $\epsilon_{\text{epi}} = 2.0$ px), where $\mathbf{u}, \mathbf{u}'$ are homogeneous pixel coordinates. This geometric constraint reduces false matches caused by repetitive texture patterns prevalent in endoscopic scenes.

Feature matching with cross-check.. Within the search window, correspondences are established by minimizing the ℓ_2 feature distance $k^* = \arg \min_{k \in W_r} \|\mathbf{f}(u, v) - \mathbf{f}(u_k, v_k)\|_2$ subject to the epipolar constraint. A cross-check (reverse match consistency within 3 px) rejects asymmetric matches, and a depth consistency filter removes
295 points with depth exceeding $5\times$ the median depth.

Multi-frame metric fusion.. The dense correspondences from each frame pair are triangulated into 3D points using the camera intrinsics and poses, then fused across all frames.

Triangulation.. Given matched homogeneous pixel coordinates $(\mathbf{u}_i, \mathbf{u}_j)$ in frames i and j , the 3D point is obtained by linear triangulation of the two projection matrices $\mathbf{P}_i = \mathbf{K}[\mathbf{I} | \mathbf{0}]$ and $\mathbf{P}_j = \mathbf{K}[\mathbf{R} | \mathbf{t}]$:

$$\tilde{\mathbf{X}}_k = \arg \min_{\tilde{\mathbf{X}}} \|\mathbf{A}\tilde{\mathbf{X}}\|^2, \quad \mathbf{A} = \begin{bmatrix} [\mathbf{u}_i]_{\times} \mathbf{P}_i \\ [\mathbf{u}_j]_{\times} \mathbf{P}_j \end{bmatrix}, \quad (8)$$

solved via SVD, where $\tilde{\mathbf{X}}$ is homogeneous, $[\mathbf{u}]_{\times}$ is the 3×3 skew-symmetric matrix of $\mathbf{u}$, and the metric point is $\mathbf{X}_k$ after dehomogenisation. Points with reprojection error exceeding $\tau_{\text{reproj}} = 2.0$ px in either view are discarded.

Multi-view consistency.. A triangulated 3D point $\mathbf{X}_k$ is retained only if it is consistent in at least two of its observing views: $\frac{1}{|\mathcal{V}|} \sum_{m \in \mathcal{V}} \mathbb{I}[\|\pi(\mathbf{K}\mathbf{T}_m\mathbf{X}_k) - \mathbf{u}_{m,k}\| < \tau_{\text{reproj}}] \geq \frac{2}{|\mathcal{V}|}$, where $\mathcal{V}$ is the set of views observing $\mathbf{X}_k$, $\mathbf{T}_m$ is the pose of view m , $\mathbf{u}_{m,k}$ is the observed pixel, $\pi(\cdot)$ is perspective projection, and $\mathbb{I}[\cdot]$ is the indicator function. This multi-view check removes spurious outliers from specular highlights.

Voxel-grid deduplication.. Overlapping reconstructions from adjacent frame pairs produce duplicate points. We apply voxel-grid downsampling (resolution: 1.0 mm) that averages positions and colors within each voxel, followed by statistical outlier removal ($k=20$ neighbors, 2.0σ threshold), yielding the final fused point cloud $\mathcal{P}$.

Complexity analysis.. The per-frame computational cost is dominated by ICP registration ($O(|\mathcal{V}| \log |\mathcal{V}|)$ per iteration with a KD-tree) and dense matching ($O(K_{\text{samp}} \cdot r^2)$, where $K_{\text{samp}} = HW/c_{\text{cell}}^2$ is the number of sampled grid points and r is the search radius). With the default parameters, the pipeline processes each frame pair in approximately 1.2 s on a single CPU core, making it suitable for offline surgical analysis.

3.8. Calibrated End-to-End Stereo Variant

The classical pipeline of Sections 3.3–3.7 derives its heightfield from image gradients and estimates pose through ICP on these pseudo-3D surfaces. The gradient-based

heightfield is a *proxy* for true scene geometry: it encodes edge structure rather than
 325 metric depth, and its accuracy is bounded by the gradient signal in texture-poor tissue.
 This section describes an optional extension used only when stereo imaging hardware
 is available. A learned disparity estimate replaces the gradient-derived height, and
 the heightfield then refines rather than replaces that metric depth. Depth results are
 reported in Table 6.

330 *Problem formulation..* Given a rectified stereo pair (I_l, I_r) with known intrinsics $\mathbf{K}$
 (focal length f_x along the disparity direction) and baseline b , the metric depth $d(u, v)$
 at pixel (u, v) relates to the disparity $d_{\text{disp}}(u, v)$ by the pinhole triangulation:

$$d(u, v) = \frac{f_x b}{d_{\text{disp}}(u, v)}. \quad (9)$$

The key challenge is estimating a dense, sub-pixel disparity field that is robust to the
 weak texture, specular reflections, and non-Lambertian surfaces typical of endoscopic
 335 tissue. We address this with a differentiable stereo network whose disparity output
 is (i) converted to metric depth via (9), (ii) lifted to a heightfield that inherits the
 geometric structure of the classical pipeline, and (iii) used jointly with the heightfield
 for pose estimation. The three modules are trained end-to-end so that the disparity
 head is optimised not only for depth accuracy but also for the downstream geometric
 340 consistency that the heightfield and pose modules require.

Encoder with attention.. Each image is encoded by a ResNet-based backbone with
 BasicBlock stages producing multi-scale features $\{\Phi^{(s)}\}_{s=1}^S$ at S resolutions (written Φ
 rather than $\mathbf{F}$ to avoid collision with the fundamental matrix). A Convolutional Block
 Attention Module (CBAM) [31] refines each stage, applying channel attention followed
 345 by spatial attention:

$$\Phi'^{(s)} = \text{CA}(\Phi^{(s)}) \otimes \Phi^{(s)}, \quad \Phi''^{(s)} = \text{SA}(\Phi'^{(s)}) \otimes \Phi'^{(s)}, \quad (10)$$

where $\text{CA}(\cdot) = \sigma(\text{MLP}(\text{GAP}(\cdot) + \text{GMP}(\cdot)))$ pools across spatial dimensions to re-
 weight channels (GAP/GMP: global average/max pooling, σ : sigmoid, MLP: shared
 two-layer perceptron), and $\text{SA}(\cdot) = \sigma(f^{7 \times 7}([\text{GAP}; \text{GMP}]))$ pools across channels to
 re-weight spatial locations ($f^{7 \times 7}$: 7×7 convolution, $[\cdot; \cdot]$: channel concatenation).

350 Channel attention suppresses uninformative feature channels (e.g., those dominated by specular highlights), while spatial attention focuses on structurally informative regions, both important in endoscopic imagery where illumination is highly non-uniform. The encoded features are aggregated to a single feature volume Φ_l, Φ_r at the highest resolution.

355 *Differentiable cost-volume aggregation.* Stereo cost aggregation constructs a 3D cost volume $\mathbf{C} \in \mathbb{R}^{D \times H \times W}$ by shifting the right feature map across D disparity hypotheses and computing a per-pixel, per-disparity matching cost as the channel-wise mean of the element-wise product of the left and shifted-right feature vectors:

$$\mathbf{C}(d, u, v) = \frac{1}{C} \sum_{c=1}^C \Phi_l^{(c)}(u, v) \Phi_r^{(c)}(u - d, v), \quad (11)$$

where C is the feature-channel count and $\Phi^{(c)}$ denotes the c -th channel; the variant V4
360 instead forms the cost from group-wise correlation [32]. The raw cost volume is regularised by a cascade of 3D convolutions [33] operating over the joint spatial–disparity space, which enforces disparity smoothness while preserving depth discontinuities at tissue boundaries. This 3D aggregation is essential because 2D convolutions cannot propagate evidence along the disparity axis and therefore produce fuzzy edges in
365 low-texture regions. A final regressor converts the regularised volume to a sub-pixel disparity map $\hat{d}_{\text{disp}}$.

Heightfield head and joint pose. The predicted disparity $\hat{d}_{\text{disp}}$ is converted to metric depth $\hat{d}(u, v)$ via (9) and lifted to a heightfield $\hat{\mathcal{H}}_t$ using the same grid construction as Section 3.3, with the learned depth replacing the gradient proxy via $\hat{h}(u, v) =$
370 $z_s \cdot (\hat{d}(u, v) - \min \hat{d}) / (\max \hat{d} - \min \hat{d})$. A pose-estimation head consumes the left–right features jointly with the heightfield to predict the relative transformation $\mathbf{T}_{t+1,t}$, replacing the ICP refinement when the network is used end-to-end. The three losses are combined: $\mathcal{L} = \lambda_d \mathcal{L}_{\text{disp}} + \lambda_h \mathcal{L}_{\text{height}} + \lambda_p \mathcal{L}_{\text{pose}}$, where $\lambda_d, \lambda_h, \lambda_p$ are scalar loss weights (default $\lambda_d=1.0, \lambda_h=0.1, \lambda_p=0.5$), $\mathcal{L}_{\text{disp}}$ is a smooth- L_1 disparity loss against stereo ground truth, $\mathcal{L}_{\text{height}}$ is an L_2 heightfield (shape) loss against the pseudo-3D target, and
375 $\mathcal{L}_{\text{pose}}$ is a geodesic pose loss [34]. Joint optimisation couples disparity accuracy with

geometric consistency: the heightfield and pose losses regularise the disparity toward configurations that yield well-conditioned 3D structure, which is exactly the property the classical pipeline relies upon for ICP convergence. Architectural variants V2–
380 V4 differ in the encoder (custom vs. pretrained ResNet) and cost-volume construction (channel-wise dot-product vs. group-wise correlation [32]); their detailed ablation is reported in Supplementary Material, Section E.

Relationship to the classical pipeline.. The two formulations are complementary. The classical pipeline (Sections 3.3–3.7) is calibration-free and operates from monocular
385 gradients, making it applicable to legacy footage and uncalibrated sequences, but its heightfield is a structural proxy rather than metric depth. The end-to-end variant requires a calibrated stereo pair and a training set, but yields metric depth whose accuracy is no longer bounded by gradient quality. The pseudo-3D heightfield is retained in both: in the classical case it is the primary geometric signal, and in the learned case it
390 is the interface through which the disparity is regularised for 3D consistency. This shared heightfield representation is what allows the ICP and dense-matching modules of Section 3.4 to be reused without modification.

4. Experiments

4.1. Dataset and Evaluation Protocol

395 *Dataset..* We evaluate all methods on an endoscopic chessboard sequence of 100 frames captured with a Karl Storz laparoscopic camera at a native resolution of 1280×720 , under endoscopic imaging conditions including specular reflections, partial occlusion by surgical instruments, and non-rigid tissue deformation. For our classical pipeline the frames are downsampled by a factor of two to 640×358 to bound the per-frame
400 computational cost, and the camera matrix used internally is the image-diagonal estimate $f_x = f_y = \sqrt{640^2 + 358^2} \approx 733.3$ px with the principal point at the image centre; the depth and pose evaluation, however, is performed at the native 1280×720 resolution using the calibrated intrinsics ($f_x = f_y = 2304$ px, principal point at image centre), so that the ground-truth corner depths and the `solvePnP` pose reference are as accurate
405 as possible. The chessboard marker has 8×5 inner corners (40 points) with square

side length $s_{sq} = 5.0$ mm, providing the independent pose reference via `solvePnP` (reprojection RMSE < 0.5 px) used only for evaluation, not as an input to our pipeline.

Input protocol. Each method takes the same 50 frames (the first 50 frames of the 100-frame sequence) as input, producing: (1) a 3D point cloud representing the reconstructed scene; and (2) estimated camera poses; for our method the poses are estimated internally
410 by the heightfield-ICP pipeline of Section 3.4 (not supplied as ground truth), whereas the baselines estimate poses with their own front-ends.

Evaluation metrics. We adopt six complementary metrics spanning geometric quality and pose accuracy:

- 415 • **PCA shape ratio** (λ_3/λ_1): Ratio of the smallest to largest eigenvalue of the point cloud’s covariance matrix $\mathbf{C} = \frac{1}{K} \sum_k (\mathbf{X}_k - \bar{\mathbf{X}})(\mathbf{X}_k - \bar{\mathbf{X}})^\top$. A well-distributed 3D structure satisfies $\lambda_3/\lambda_1 \geq 0.25$; values below indicate degenerate geometry.
- **Shape classification**: Normal 3D ($\lambda_3/\lambda_1 \geq 0.25$), Flat ($0.15 \leq \lambda_3/\lambda_1 < 0.25$), Cone ($\lambda_3/\lambda_1 < 0.15$). This threshold is chosen so that a Normal 3D point cloud
420 has comparable extent along all three principal axes, which is the minimum requirement for representing a 3D surgical scene.
- **ATE** (Absolute Trajectory Error) [35]: RMSE of camera positions after Sim3 alignment (scale, rotation, translation) between estimated and ground-truth trajectories.
- 425 • **RPE** (Relative Pose Error): Mean relative translation error between consecutive frame pairs, measuring local consistency independent of global drift. A global rotation error is not reported: the benchmark sequence is translation-dominant (reference relative rotation $\approx 0.3^\circ$ per frame), so the rotational part of a position-fitted Sim3 alignment is ill-conditioned.
- 430 • **Sim3 scale factor**: The scale component of the Sim3 alignment; $1.0\times$ indicates correct metric scale.
- **Depth metrics**: absolute relative error (AbsRel), root-mean-square error (RMSE), and threshold accuracy (δ_1), computed at chessboard corner positions after affine alignment ($gt = a \cdot pred + b$), measuring per-point depth accuracy

435 against PnP-derived ground truth. The affine (scale+shift) alignment is used
 here, rather than the median-scale alignment adopted on the public stereo datasets
 (Sections 4.4–4.7), because the chessboard ground truth is PnP-derived and the
 heightfield/SGM disparity is a structural proxy whose systematic offset must be
 removed before the per-corner error is meaningful; the chessboard depth figures
 440 are therefore not comparable in magnitude with the other datasets, only within
 this table.

4.2. Compared Methods

We compare against eight representative baselines spanning classical SfM, feed-forward 3D reconstruction networks, and pose regression:

- 445 • **Stereo SGM**: Semi-global matching (SGM) [16] stereo pipeline with the same
 heightfield-ICP poses as Ours, serving as a classical stereo baseline with identical
 pose input.
- **ORB-SfM** [6]: ORB feature extraction [18] with 5-point essential matrix esti-
 mation [36] and incremental pose chaining via `recoverPose`. Triangulation via
 450 `triangulatePoints`. No bundle adjustment.
- **COLMAP** [7]: SIFT features [19] with incremental SfM and full bundle adjust-
 ment. A well-established classical baseline.
- **VGGT** [5] (CVPR 2025 Best Paper): Feed-forward Vision Geometry Grounded
 Transformer that jointly predicts camera poses, depth maps, and 3D point clouds
 455 from an image sequence in a single forward pass. DINOv2 [37] backbone with
 1.2B parameters.
- **Reloc3R** [24] (CVPR 2025): Relative pose regression network with motion
 averaging for global pose estimation. Extends DUST3R for camera relocalization.
 Outputs poses only (no point cloud).
- 460 • **DUST3R** [3] (CVPR 2024): Dense uncalibrated stereo 3D reconstruction via
 paired pointmap prediction and global alignment optimization (Levenberg–
 Marquardt).
- **MASt3R** [4] (ECCV 2024): Extension of DUST3R with local feature matching

and metric pointmaps for improved 3D reconstruction and camera pose estimation.

- **MUS_t3R** [23] (CVPR 2025): Multi-view extension of DUS_t3R that processes $N > 2$ views simultaneously with a memory network, producing globally consistent pointmaps and poses without pairwise alignment.

Implementation details. All methods use the same 50 input frames (the first 50 frames of the 100-frame sequence) with default hyperparameters and their officially released pretrained models; each method resizes the frames to its own operating resolution (our pipeline to 640×358 , MUS_t3R to 224×224 , and the other feed-forward networks to their native input sizes), so the comparison is on the same image content rather than on identical pixel grids. No method receives camera intrinsics as input (except ORB-SfM, which uses the chessboard-calibrated intrinsics for `recoverPose`). Ours estimates camera poses internally by registering consecutive heightfield meshes with FPFH features, RANSAC, and point-to-point ICP (Section 3.4); the chessboard is used only to obtain intrinsics and an independent pose reference for evaluation. DUS_t3R and MAS_t3R are run with the global alignment optimizer (pairwise then global refinement). VGGT processes 50 frames with sub-sequence windowing (8 frames per window). MUS_t3R uses a memory network with 8 keyframes at 224×224 resolution. For our method, we use the following parameters: gradient threshold $\tau_{\%} = 68$, mesh cell size $c_{\text{cell}} = 3$, coarse alignment via PCA principal-axis alignment followed (on failure) by FPFH-feature-matched RANSAC with a 5 $\times$ -enlarged correspondence distance, refinement via point-to-point ICP with an adaptive (scale-derived) correspondence distance, dense matching search radius $r=25$, and epipolar band $\epsilon_{\text{epi}} = 2.0$. The evaluation uses a fixed random seed (42) for reproducibility of the PCA sub-sampling step. Complete hyperparameter and hardware details are provided in the Supplementary Material, Section A.

4.3. Quantitative Results

Table 2 presents the comprehensive quantitative comparison on the benchmark, covering depth accuracy, 3D shape quality, and camera pose precision across nine

Table 2: Comparison on the endoscopic chessboard benchmark (100 frames captured at 1280×720, 50 used as input; our pipeline runs on 640×358 downsampled frames, Karl Storz laparoscope). Depth metrics (AbsRel, RMSE, δ_1) computed at chessboard corners (40 per frame, 4000 total, ground-truth depth range 225–784 mm) after affine alignment ($gt = a \cdot pred + b$); the RMSE values are therefore in millimetres against depths of several hundred millimetres, so an RMSE of, e.g., 109 mm corresponds to a relative error of about 22%, consistent with the reported AbsRel. λ_3/λ_1 : 3D shape ratio (≥ 0.25 Normal 3D; 0.15–0.25 Flat; < 0.15 Cone). ATE (mm) and Sim3 scale are computed against the independent solvePnP chessboard reference trajectory; all methods are aligned to the same reference for comparability. [†] Reloc3R produces camera poses only and does not output a point cloud; consequently its depth-accuracy (λ_3/λ_1 , AbsRel, RMSE, δ_1) and point-count columns are reported as N/A, while pose metrics (ATE, Scale) remain available. [‡] Ours and Stereo SGM share the same heightfield-ICP trajectory, so their pose metrics are identical. MUST3R is evaluated on 20 of the 50 frames (its multi-view memory processes a fixed keyframe budget). **Bold**: best per metric.

Method	Pts	Depth Accuracy			3D Shape		Pose	
		AbsRel↓	RMSE↓	δ_1 ↑	λ_3/λ_1	Shape	ATE↓	Scale
Ours [‡]	11,735	0.2465	109.19	0.5723	0.4205	Normal 3D	3.89	1.11
Stereo SGM [‡]	6,917	0.2126	86.86	0.6414	0.3123	Normal 3D	3.89	1.11
VGGT [5]	324,162	0.0091	3.74	1.0000	0.1938	Flat	3.02	90.89
DUST3R [3]	229,069	0.1833	64.65	0.6860	0.0865	Cone	30.78	647.07
MASt3R [4]	314,474	0.1843	76.62	0.6265	0.2201	Flat	36.86	216.65
MUST3R [23]	1,003,520	0.0853	24.06	1.0000	0.0647	Cone	12.13	9.85
COLMAP [7]	1,994	0.1692	69.83	0.7521	0.1312	Cone	20.10	8.86
ORB-SfM [6]	6,114	0.1779	69.32	0.6627	0.0115	Cone	3.79	3.71
Reloc3R [24] [†]	0	N/A	N/A	N/A	N/A	—	26.84	119.91




Figure 2: Chessboard shape, one method per row. The input row is the board cropped from four frames. The Ours row is the corresponding RGB heightfield (green). Each later row is that method’s fused cloud in four views, drawn the same way. The bar is λ_3/λ_1 . The dashed line is the Normal 3D threshold 0.25. Only Ours (0.42) and Stereo SGM (0.31) clear it. SGM has no fused chessboard cloud in this figure; its ratio is the table value. Colour is native RGB. A depth pseudocolour would paint the Ours ridges as a mottled blob and a collapsed slab as a rectangle.

methods. The qualitative evidence is split across three figures so that each page stays a small raster rather than one vector point-cloud plate: chessboard shape (Fig. 2), measured trajectories on twelve of the seventeen sequences (Fig. 3), and clinical surfaces (Fig. 4).

Point cloud shape quality. Fig. 2 colours every chessboard cloud by its native RGB, not by height along the viewing axis. That choice is deliberate. A Z-height pseudocolour maps the third principal component onto hue: because the proposed heightfield is a gradient-ridge surface with genuine thickness, the same colouring turns the Ours face-on view into an unstructured mottled blob, while VGGT and MAST3R, having collapsed into thin planes, project to tidy rectangles that look more like a chessboard than the volumetric reconstruction does. The rectangle is the silhouette of a collapsed slab, not evidence that the checker geometry was recovered. The top rows of Fig. 2 place the




Figure 3: Measured trajectories on twelve real sequences, each panel the same size, projected onto the plane of the ground-truth motion. Legend: black, ground truth; green, Ours-FullGauge; red, ORB-SfM. The green title is that panel's FullGauge ATE in millimetres. The bottom row is Sim3 ATE on all seventeen sequences. Ours is FullGauge on real sequences and multi-stride BA on the four simulations. VGGT-1B is blue and ORB-SfM is red.




Figure 4: Clinical surfaces in equal panels. Rows are EndoNeRF, SCARED, and Stereo-Lap. Each method is shown face-on and from the side: Ours (green), Stereo SGM, and Mono. Ours keeps a textured surface with SGM. Mono collapses to a cone. The bottom row is six views of the dense EndoNeRF end-to-end cloud. That file has no RGB, so colour is thickness.

505 cropped board next to the Ours heightfield and, on the rows below, the same four-view layout for VGGT, MAST3R, COLMAP, DUST3R, MUST3R, and ORB-SfM. The bar is the shape comparison: only Ours ($\lambda_3/\lambda_1 = 0.42$) and Stereo SGM (0.31) lie above the Normal 3D line at 0.25. VGGT (0.19) and MAST3R (0.22) are Flat; COLMAP (0.13), DUST3R (0.09), MUST3R (0.06), and ORB-SfM (0.01) are Cone. This figure

510 uses only chessboard reconstructions. Clinical EndoNeRF tissue point clouds appear exclusively in Fig. 4 and are not inserted into the chessboard columns in place of Stereo SGM; the SGM chessboard shape ratio ($\lambda_3/\lambda_1 = 0.31$, also Normal 3D) is the table entry in Table 2, not a tissue cloud reused as a board. This degeneration is *independent of scale*: the PCA eigenvalue ratio measures intrinsic volumetric spread, and a flat

515 or cone-like point cloud cannot represent a 3D surgical cavity regardless of post-hoc rescaling. Fig. 5 visualizes this separation across all datasets, where only our method and SGM enter the Normal 3D domain while all feed-forward baselines remain trapped in Flat or Cone regimes.

Trajectory tracking fidelity. Fig. 3 draws twelve real tracks on the plane of the ground-truth motion, in equal panels. The green curve sits on the black reference; the red ORB-SfM curve does not. The bottom bar is the Sim3 ATE of all seventeen sequences. On the test recording `traj_023422`, where the endoscope executes large lateral and axial excursions with working distances up to 260 mm, full metric scale gauge injection (chessboard PnP) locks onto the ground truth at 4.55 mm ATE, against VGGT-1B at 16.86 mm and ORB-SfM at 68.59 mm. The same FullGauge lock holds on the other real panels (for example 1.31 mm on `traj_015918` and 2.05 mm on `rigid_142902`). On the pattern-free cavity `s10097`, multi-stride bundle adjustment reaches 24.86 mm ATE and beats VGGT-1B (37.25 mm) by 12.39 mm. Fig. 4 uses one panel size for EndoNeRF, SCARED, and Stereo-Lap, with a face view and a side view of Ours, Stereo SGM, and Mono. The green frame is Ours: a textured surface, where Mono falls into a cone. The bottom row is six views of the dense EndoNeRF end-to-end reconstruction.

Depth accuracy at chessboard corners. To evaluate metric depth, we generate ground-truth depths at the 4000 chessboard corners (40 per frame $\times$ 100 frames) via PnP with known board geometry (reprojection RMSE < 0.5 px). After affine alignment ($gt = a \cdot pred + b$), VGGT achieves the best depth accuracy (AbsRel=0.009, $\delta_1=1.0$), demonstrating that feedforward networks can estimate accurate relative depth when the scene is well-textured. MUST3R ranks second (AbsRel=0.085, $\delta_1=1.0$), benefiting from its multi-view extension of DUST3R. Our method and Stereo SGM achieve AbsRel of 0.247 and 0.213, respectively, higher (worse) than COLMAP (0.169) and ORB-SfM (0.178), and substantially higher than the feed-forward baselines VGGT (0.009) and MUST3R (0.085). This is expected: our heightfield-based reconstruction prioritizes 3D shape preservation over per-pixel depth precision, and the SGM disparity at chessboard corners is limited by the sparse gradient structure of the calibration pattern. The method therefore trades per-corner depth accuracy for a well-formed 3D shape; the two quantities answer different questions and should not be collapsed into a single ranking.

Square-size measurement. Affine-aligned AbsRel removes a per-corner scale and shift before the error is computed, so it is not a measurement of the 5.0 mm square in the

reconstruction’s own units. We therefore back-project the published nearest-neighbour
550 depth (25 px, the intrinsics and pose convention of the depth table) along each detected
corner ray and compare adjacent inner-corner lengths with 5.0 mm. No affine and no
Sim3 are applied. The same rays recover the square from the PnP depths (mean 5.00 mm,
MAE 0.047 mm), so the edge geometry is not the source of the error. On 1874 edges our
native median length is 5.49 mm and the mean absolute error is 10.9 mm (interquartile
555 range 1.10–16.5 mm). ORB-SfM has median 6.26 mm and MAE 5.89 mm. VGGT
and COLMAP have medians 0.019 mm and 0.072 mm. An affine fit on these same
matches recovers AbsRel 0.248 for our cloud, against the published 0.247. The median
near 5 mm is not a millimetre gauge: the absolute error is larger than the square itself.
Table 3 records the comparison.

560 *Where the measurement fails..* The gap is not a missing global scale. A single similarity
fit on the first 15 frames, evaluated on the remaining frames, changes the native edge
MAE from 9.5 mm to 13.8 mm. The cloud is not a rigid copy of the board. Consecutive-
frame triangulation with the stored poses and $\mathbf{K}$ makes the cause specific: within one
pair the edges are consistent, but the pair-median length wanders (interquartile range
565 0.16–0.56 in the pose units). The heightfield rigid motion and the pinhole triangulation
do not share a unit of translation. We therefore keep the heightfield for registration, and
apply a sparse baseline gauge only at triangulation: every tenth pair sets one scale so
that its median reconstructed edge equals 5.0 mm, and the scale is interpolated to the
pairs in between. On the 77 held-out pairs the median edge is 5.03 mm and the median
570 per-pair MAE is 1.45 mm; 66 of 77 pairs have MAE below 3 mm. Two pairs (starting
at frames 5 and 13) remain failures, so the mean MAE is 8.2 mm. The same gauge on
the other methods’ poses gives a higher hold-out median MAE: COLMAP 3.41 mm (7
of 25 pairs below 3 mm), VGGT 4.65 mm (1 of 34), and ORB-SfM 4.85 mm (0 of 21).
This is the measurement on which the heightfield poses beat the compared methods. It
575 is not applied to the fused cloud in Table 3; the correspondences here are the detected
grid, not the 7-D tissue matches.

Photometric assist inside the pose window.. The 7-D feature keeps one grey value,
so a periodic pattern is ambiguous: adjacent chessboard corners are 45 px apart, and

Table 3: Native adjacent-corner length on the chessboard (5.0 mm squares). Depth at each corner is the nearest fused point within 25 px, back-projected along the detected ray. No affine alignment and no Sim3. The ground-truth row uses PnP depths on the same rays.

Method	Edges	Median (mm)	MAE (mm)	IQR (mm)
PnP depths	6700	5.00	0.047	4.96–5.04
Ours	1874	5.49	10.9	1.10–16.5
ORB-SfM	2010	6.26	5.89	1.83–12.4
COLMAP	1825	0.072	4.89	0.048–0.12
VGGT	2531	0.019	4.98	0.019–0.020

mapping $(u, v, 0)$ through the heightfield pose lands a median of 43 px from the true
580 corner. The fundamental matrix built from that same pose misses the true corner by
a median of 621 px, so the search is not run along that line. Every tenth pair instead
records the pixel bias of the predictor. On the other pairs the bias is interpolated and
ZNCC (17×17 template, normalised correlation) is maximised inside an 18 px window,
under half a board period, so the match cannot jump to the next square. The detector is
585 not used on these pairs. Across 75 held-out pairs the median match error is 0.73 px (68%
within 3 px). With the same interpolated baseline scale, the median edge is 4.97 mm
and the median per-pair MAE is 2.01 mm; 55 of 75 pairs are below 3 mm. The mean
MAE is 17.4 mm because the remaining pairs still slip. A learned depth order, such as
VGGT, does not see this failure: the two squares lie on one plane, so their depths agree.
590 VGGT remains a veto for matches that jump off the surface, not a replacement for the
heightfield cloud.

On the role of the heightfield-derived poses.. A natural concern is whether the Nor-
mal 3D shape advantage stems from the heightfield representation or simply from
accurate camera poses. We clarify two points that address this concern directly. *First*,
595 the poses used by our pipeline on the chessboard benchmark are not externally supplied
ground truth: they are estimated *by the heightfield pipeline itself*, by registering consec-
utive 2.5-D heightfield meshes with FPFH features, RANSAC, and point-to-point ICP

Table 4: Pose-source ablation on the chessboard benchmark (50 frames). Both rows use the identical dense-matching and fusion pipeline; only the camera-pose source differs. The heightfield-ICP poses are estimated end-to-end by the proposed pipeline; the essential-matrix poses use a conventional 2D SIFT front-end. ATE and the Sim3 scale factor are computed against the independent `solvePnP` chessboard reference. The point-cloud shape ratio λ_3/λ_1 is reference-frame invariant (verified by re-evaluating in the chessboard frame), so the Cone degeneration under the 2D front-end is a genuine geometric collapse, not a coordinate artefact.

Pose source	Points	λ_3/λ_1	Shape	ATE↓	Sim3 scale
Heightfield-ICP (Ours)	11,735	0.4205	Normal 3D	3.89	1.11
Essential-matrix (2D)	8,283	0.0190	Cone	14.70	6.51

(Section 3.4) and chaining the resulting relative transforms into absolute camera poses. The Normal 3D result ($\lambda_3/\lambda_1 = 0.42$) is therefore obtained end-to-end from monocular
600 video, with no separate pose sensor. *Second*, to confirm that the heightfield geometry (rather than pose accuracy alone) drives the result, we replaced these heightfield-ICP poses with poses estimated from a conventional 2D front-end (SIFT feature matching, essential-matrix RANSAC, and `recoverPose`), while keeping every other component of the pipeline fixed. Under this weaker pose source the reconstruction degenerates to
605 a Cone with $\lambda_3/\lambda_1 \approx 0.02$, indistinguishable from the learned and SfM baselines. The implication is that the 3D geometric signal carried by the heightfield is what makes both the pose estimation and the downstream reconstruction well-conditioned: a conventional 2D pose front-end, which lacks this 3D signal, cannot sustain the Normal 3D geometry even when feeding the same dense-matching and fusion modules. The height-
610 field thus serves a dual role: as the substrate for pose estimation and as the representation fused into the final point cloud. This is the property the baselines, relying on either 2D correspondences or learned natural-image priors, do not share. Table 4 summarises this pose-source ablation.

Pose accuracy and scale drift. Against the independent `solvePnP` reference, Ours
615 attains an ATE of 3.89 mm, within 0.9 mm of the best feed-forward result (VGGT, 3.02 mm) and comparable to ORB-SfM (3.79 mm) (Fig. 6(b)). The decisive difference is metric scale (Fig. 6(c)): Ours recovers a near-ideal Sim3 scale of 1.11 $\times$, whereas every

baseline drifts—ORB-SfM 3.7×, COLMAP 8.9×, MUST3R 9.9×, VGGT 91×, Reloc3R 120×, MAST3R 217×, and DUST3R 647×. Such scale ambiguity precludes clinical
620 metric measurements such as instrument-to-tissue distance or lesion size estimation, even when the trajectory shape itself is accurate. Because the sequence is translation-dominant (the reference relative rotation is only 0.3° per frame), the rotational part of a position-fitted Sim3 alignment is ill-conditioned; we therefore report relative translation error instead of a global rotation error. Per-frame relative translation error is comparable
625 across all methods (2.2–3.3 mm; Ours 3.10 mm), indicating that no method accumulates significant local drift on this benchmark (Fig. 6(d)).

Density vs. geometric quality. Point cloud density does not correlate with 3D quality: MUST3R (1.0M points), VGGT (324K), and MAST3R (314K) produce two orders of magnitude more points than Ours (11.7K), yet their λ_3/λ_1 ratios are less than half.
630 Feedforward networks thus generate dense but geometrically unreliable reconstructions in surgical scenes, where repetitive texture and specular reflections violate their learned matching priors. Classical methods (ORB-SfM, COLMAP) fail due to insufficient feature correspondences; learned methods fail geometrically despite producing dense outputs. Our method addresses both failure modes by exploiting gradient-based geometric anchors for reconstruction, and uses the same heightfield geometry to estimate
635 the camera poses that the baselines recover from 2D correspondences or learned priors.

4.4. Evaluation on Stereo Endoscopic Data

To evaluate on a public benchmark without calibration-pattern anchors, we use the EndoNeRF dataset [38], which provides stereo endoscopic video from a DaVinci
640 surgical robot with ground-truth depth from the right-view stereo camera. We evaluate on two sequences (50 frames each) with intrinsics from LLFF poses ($f = 569.47$ px, 640×512). We compare five methods: Stereo SGM, Mono (gradient-inverse depth), ORB-SfM, MAST3R [4], and Ours (using the dataset-provided LLFF poses). We note an asymmetry in pose inputs on this public benchmark: Ours and Stereo SGM use
645 the dataset-provided poses, whereas ORB-SfM and MAST3R estimate their own. This is unavoidable here because our classical pipeline does not yet include a pose front-end tuned for narrow-baseline surgical footage (see Section 5, Limitations), and we

PCA Shape Distribution across all datasets




Figure 5: PCA shape distribution (λ_2/λ_1 vs. λ_3/λ_1) across all datasets. The Normal 3D region ($\lambda_3/\lambda_1 \geq 0.25$, green), Flat ($0.15 - 0.25$, orange), and Cone (< 0.15 , red) boundaries are marked. Large markers: chessboard benchmark; small markers: EndoNeRF, SCARED, and Stereo Lap sequences. Only Ours (star) and a subset of Ours/SGM points on other datasets attain Normal 3D; all six baselines fall within Flat or Cone regions.




Figure 6: Quantitative comparison across all methods on the endoscopic chessboard benchmark. (a) PCA shape ratio λ_3/λ_1 : Ours (0.42) and Stereo SGM (0.31) exceed the Normal 3D threshold (0.25, green dashed line). (b) ATE after Sim(3) alignment against the independent `solvePnP` chessboard reference. (c) Sim(3) scale factor (logarithmic scale): the learned and SfM baselines drift by $3.7\times$ – $647\times$, whereas Ours recovers a near-ideal $1.11\times$ (ideal value $1.0\times$, green dashed line). (d) Relative translation error per frame pair: all methods remain within 2.2–3.3 mm, indicating comparable local consistency. Reloc3R outputs no point cloud and is omitted from (a); Stereo SGM shares the heightfield-ICP trajectory of Ours and is omitted from (b)–(d) to avoid double counting.

therefore restrict the EndoNeRF comparison to depth accuracy (where, as reported below, Ours matches Stereo SGM by construction and we claim no advantage over it) rather than to pose accuracy. Depth metrics (AbsRel, RMSE, $\delta < 1.25$) are computed after median-scale alignment against stereo ground truth.

Ground-truth representations and cross-dataset comparability. The three depth-evaluation datasets used in this paper store ground truth in heterogeneous representations, and the corresponding AbsRel/RMSE values are therefore *not* directly comparable across datasets. (i) EndoNeRF provides depth as 8-bit PNG maps (value range $\approx 0-93$), i.e. a scale-ambiguous relative depth rather than an absolute metric quantity; consequently, all EndoNeRF metrics are computed after per-frame median alignment of each prediction to the ground truth, and the resulting AbsRel (e.g., 0.0285) is a scale-aligned relative error, not a metric error. (ii) SCARED provides float32 structured-light depth maps in millimetres (median ≈ 22 mm) with NaN values marking invalid regions; metrics are computed over the finite, positive overlap of prediction and ground truth, and the larger AbsRel values ($\approx 2.5-2.8$) reflect both the metric (millimetre) scale and the zero-shot domain transfer from the EndoNeRF-trained model. (iii) The Stereo-Laparoscopic dataset provides metric depth in millimetres. Because of these scale and representation differences, depth-accuracy figures should be interpreted *within* each dataset and not compared numerically across datasets; cross-dataset conclusions are drawn from within-dataset method rankings and from the qualitative visualisations in Fig. 12. The 3D shape statistic (λ_3/λ_1 , Table 11) is scale-invariant and is the only quantity directly comparable across datasets.

Table 5 presents the results.

Depth accuracy. Ours and Stereo SGM achieve identical depth accuracy (AbsRel $\approx 0.031-0.034$), because the depth source is the same SGM disparity and the heightfield only adds a confidence map; we therefore claim no depth-accuracy gain over SGM here. Both substantially outperform the monocular baseline (AbsRel > 0.4).

3D shape quality. On *Pulling Soft Tissues*, both Ours and SGM achieve Normal 3D ($\lambda_3/\lambda_1 = 0.29$ and 0.32 , respectively), while MAST3R ($\lambda_3/\lambda_1 = 0.13$) produces Cone

Table 5: Evaluation on the EndoNeRF stereo endoscopic dataset [38] (DaVinci surgical robot, 640×512, 50 frames per sequence). Depth metrics (AbsRel, RMSE, δ_1) computed after median-scale alignment against right-view stereo ground truth. λ_3/λ_1 : 3D shape ratio (≥ 0.25 Normal 3D; 0.15–0.25 Flat; < 0.15 Cone). [†] Depth evaluation infeasible due to severe scale/frame misalignment. **Bold**: best per sequence per metric.

Seq	Method	AbsRel↓	RMSE↓	δ_1 ↑	Shape	λ_3/λ_1
Pulling	Mono	0.446	53.0	0.333	Flat	0.21
	MASt3R [4]		†		Cone	0.13
	Stereo SGM	0.031	7.30	0.991	Normal 3D	0.32
	Ours	0.031	7.30	0.991	Normal 3D	0.29
Cutting	Mono	0.526	46.1	0.292	Cone	0.15
	ORB-SfM [6]		†		Cone	0.04
	Stereo SGM	0.034	3.82	0.983	Cone	0.14
	Ours	0.034	3.82	0.983	Flat	0.18

geometry. On *Cutting Tissues Twice*, Ours achieves Flat ($\lambda_3/\lambda_1 = 0.18$) while SGM degenerates to Cone ($\lambda_3/\lambda_1 = 0.14$), and ORB-SfM ($\lambda_3/\lambda_1 = 0.04$) also produces Cone geometry. The shape differences between Ours and SGM are minor and arise from stochastic point sampling in PCA computation.

Runtime.. On EndoNeRF, where dataset-provided poses are used and only the SGM disparity plus heightfield confidence map are computed (without inter-frame ICP or dense epipolar matching), Ours processes 50 frames in 3.8–4.6 s, comparable to SGM and about 100× faster than MASt3R (455 s on Pulling). The full pipeline on the chessboard benchmark, which additionally runs heightfield-ICP pose estimation and dense matching, takes 61.5 s for 50 frames (Table 12).

Summary.. On EndoNeRF the heightfield confidence map matches SGM depth accuracy while flagging unreliable disparity regions for downstream use; the depth-accuracy advantage of our framework on endoscopic tissue comes from the calibrated end-to-end stereo variant (Section 4.5), not from the heightfield-on-SGM combination.

4.5. Comparison with Depth Estimators on EndoNeRF

We also benchmark a calibrated end-to-end stereo variant of our framework against representative depth estimators on the EndoNeRF dataset [38], using the same two surgical sequences: *Pulling Soft Tissues* and *Cutting Tissues Twice*. To ensure a fair
695 comparison, all methods are evaluated on an identical frame set comprising the first 50 frames of each sequence (100 frames in total).

Compared methods. The baselines are organized into three groups. (i) Zero-shot stereo matching: FoundationStereo [39], evaluated at full scale ($s=1.0$) and half scale ($s=0.5$), and IGEV-Stereo [40] with its Middlebury [41] pretrained weights. (ii) Classical stereo
700 matching: semi-global matching (Stereo SGM) and its monocular gradient-inverse counterpart (Mono), which serve as the conventional references. (iii) Monocular foundation models: Depth Anything V2 (DA-V2) [42] in Small, Base, and Large variants. EndoNeRF ground truth is scale-ambiguous relative depth, so every method—including the proposed calibrated model—is aligned to the ground-truth depth frame-by-frame
705 via median scaling, the standard protocol for scale-ambiguous depth evaluation [13]. The stereo geometry still supplies a metric depth estimate; the alignment is used only so that errors against the 8-bit EndoNeRF maps are comparable.

Metrics. Three standard depth metrics are reported: the absolute relative error (AbsRel, $\downarrow$), the root-mean-square error (RMSE, $\downarrow$), and the threshold accuracy $\delta_{1.25}$ ($\uparrow$).
710 A common valid mask is applied across all methods, retaining ground-truth pixels satisfying $d > 0$ and $d < \max(3 p_{99}, 200)$ intersected with valid predictions, thereby suppressing invalid or saturated pixels in the EndoNeRF depth maps.

Table 6 summarizes the results. The proposed calibrated model attains the lowest overall AbsRel (0.0285) and RMSE (3.97), improving over the classical Stereo SGM by
715 11.8% and 28.6% and over the strongest zero-shot baseline, FoundationStereo, by 11.2% and 10.4%, with consistent gains on both sequences. Stereo SGM achieves a marginally higher $\delta_{1.25}$ on the pulling sequence (0.9910 vs. 0.9857), yet its overall RMSE is 40% larger (5.56 vs. 3.97): the $\delta_{1.25}$ threshold is insensitive to error magnitude and therefore masks the gross outliers present in the SGM output, which the proposed refinement
720 suppresses. All stereo methods outperform the monocular baselines by a wide margin:

Table 6: Quantitative comparison of depth estimation on EndoNeRF [38] (first 50 frames per sequence, 100 frames in total). All methods are evaluated on an *identical* frame set. The EndoNeRF ground truth is 8-bit scale-ambiguous relative depth; all metrics are therefore computed after per-frame median alignment, and the reported AbsRel/RMSE are scale-aligned relative errors (not metric errors). Scale-less baselines and the proposed calibrated model are evaluated under this same protocol. Method types: Ours = calibrated stereo + pseudo-3D; FS, IGEV = zero-shot/classical stereo matching; DA-V2, Mono = monocular. ↓/↑ denote lower/higher is better; best results are in **bold**. IGEV-Stereo is evaluated on the cutting sequence only.

Method	Cutting Tissues Twice			Pulling Soft Tissues			Overall (100 frames)		
	AbsRel↓	RMSE↓	$\delta_{1.25}$ ↑	AbsRel↓	RMSE↓	$\delta_{1.25}$ ↑	AbsRel↓	RMSE↓	$\delta_{1.25}$ ↑
Proposed	0.0252	3.31	0.9877	0.0317	4.63	0.9857	0.0285	3.97	0.9867
FoundationStereo [39]	0.0307	4.05	0.9806	0.0334	4.82	0.9817	0.0321	4.43	0.9812
FoundationStereo-0.5x [39]	0.0407	5.07	0.9757	0.0452	6.14	0.9755	0.0429	5.61	0.9756
Stereo SGM	0.0336	3.82	0.9829	0.0310	7.30	0.9910	0.0323	5.56	0.9870
IGEV-Stereo [40]	0.2254	15.17	0.5218	–	–	–	–	–	–
DA-V2-S [42]	0.1390	11.23	0.7881	0.2686	27.03	0.5088	0.2038	19.13	0.6484
DA-V2-B [42]	0.1886	13.58	0.6274	0.2861	30.62	0.5192	0.2373	22.10	0.5733
DA-V2-L [42]	0.2199	18.39	0.5810	0.2793	28.65	0.4849	0.2496	23.52	0.5330
Mono (Grad-inverse)	0.5256	46.08	0.2923	0.4455	53.04	0.3330	0.4856	49.56	0.3126




Figure 7: Depth-estimation comparison on EndoNeRF (100 frames, logarithmic scale). (a) AbsRel and (b) RMSE over the full frame set; (c) per-sequence AbsRel. The proposed calibrated model is compared against zero-shot stereo, classical stereo, and monocular baselines, and attains the lowest error on both metrics. IGEV-Stereo is evaluated on the cutting sequence only (hatched bar).

the best monocular model (DA-V2-S) reaches only AbsRel 0.204, about $7\times$ worse than the proposed model, because monocular depth is scale-ambiguous and generalizes poorly to weak texture and non-Lambertian reflections in endoscopic tissue. Among the monocular foundation models, accuracy degrades with capacity ($S > B > L$ on AbsRel); larger models encode stronger natural-image priors that bias predictions away from endoscopic appearance. IGEV-Stereo, with its Middlebury pretrained weights, attains only AbsRel 0.225 on the cutting sequence, as its disparity range and appearance priors are ill-matched to the endoscopic setting, indicating that generic stereo models require adaptation before deployment.

Fig. 7 visualizes the comparison. Panels (a) and (b) show the margin between the stereo-based and monocular methods and the consistent advantage of the proposed model on both error metrics; the logarithmic scale is used so that the stereo cluster (AbsRel ≈ 0.03) and the monocular cluster (AbsRel ≥ 0.20) remain legible together. Panel (c) shows that the pulling sequence is more challenging than the cutting sequence for every method (the proposed AbsRel rises from 0.0252 to 0.0317), yet the proposed model retains its advantage on both, indicating robustness to larger non-rigid deformation. Combining calibrated stereo geometry with the pseudo-3D refinement thus yields metric depth that is more accurate and robust than the zero-shot, classical, and monocular baselines on endoscopic data.

Table 7: Evaluation on the SCARED dataset [43] across eleven keyframes (porcine cadaver endoscopic video, structured-light depth GT). Values are averaged across keyframes within each dataset subset (DS1: 3 keyframes, DS2: 4 keyframes, DS3: 4 keyframes). Detailed per-keyframe breakdowns are reported in Supplementary Material, Section F. **Bold**: best per subset per metric.

Dataset	Method	AbsRel↓	RMSE (mm)↓	$\delta_1 \uparrow$	λ_3/λ_1	Shape
DS1 (3 KFs)	Mono	1.070	191.0	0.190	0.09	Cone
	Stereo SGM	0.079	10.82	0.957	0.10	Cone
	Ours	0.079	10.82	0.957	0.11	Cone
DS2 (4 KFs)	Mono	1.028	149.5	0.202	0.12	Cone
	Stereo SGM	0.139	19.47	0.846	0.35	Normal 3D
	Ours	0.139	19.47	0.846	0.31	Normal 3D
DS3 (4 KFs)	Mono	1.313	214.9	0.173	0.11	Cone
	Stereo SGM	0.319	34.73	0.524	0.24	Flat
	Ours	0.319	34.73	0.524	0.23	Flat

4.6. Evaluation on the SCARED Dataset

To further validate on clinical-grade stereo data, we evaluate on the SCARED dataset [43], which provides stereo endoscopic video of porcine cadavers with ground-truth camera poses and structured-light depth maps. We evaluate on three datasets (DS1: 3 keyframes, DS2: 4 keyframes, DS3: 4 keyframes); 30 frames are uniformly sampled per sequence. Camera intrinsics are obtained from the provided calibration ($f_x \approx f_y \approx 1035$ px, 1280×1024 , baseline ≈ 4.14 mm). We compare Ours (heightfield confidence-weighted SGM, using the dataset-provided poses) against Stereo SGM and Mono; all three use the same dataset-provided poses, so the comparison isolates the depth/disparity stage rather than pose estimation. Depth metrics are computed after median-scale alignment against structured-light ground truth.

Results.. Table 7 presents the summary across eleven SCARED keyframes. Across all datasets, Ours and Stereo SGM achieve *identical* depth accuracy (AbsRel, RMSE, δ_1) by construction: on SCARED our pipeline uses SGM disparity as its depth source and adds a heightfield-derived confidence map on top, so the depth figures match rather than demonstrate a depth-accuracy gain. We therefore do not claim a depth advantage

Qualitative depth comparison on the SCARED dataset (five keyframes). The proposed SGM-heightfield pipeline recovers tissue-surface depth consistent with the structured-light ground truth, whereas the monocular baseline yields structureless depth. (Turbo colormap, normalised per panel; predictions median-aligned to the GT for display.)




Figure 8: Qualitative depth comparison on the SCARED dataset (five keyframes). The proposed SGM-heightfield pipeline recovers tissue-surface depth consistent with the structured-light ground truth, whereas the monocular baseline yields structureless depth. (Turbo colormap, normalised per panel; predictions median-aligned to the GT for display.)

over SGM on this dataset; instead, the contribution on SCARED is the *confidence map* itself, which flags unreliable disparity regions (e.g., specular highlights, low-texture tissue) for downstream filtering and region-of-interest selection. Both methods substantially outperform the monocular baseline ($\text{AbsRel} > 1.0$). On 3D shape quality
760 the two methods are also comparable: both achieve Normal 3D on DS2 keyframes ($\lambda_3/\lambda_1 \approx 0.31\text{--}0.35$). DS3 keyframes exhibit higher depth error ($\text{AbsRel} > 0.3$) due to larger camera-to-scene standoff distances.

Fig. 8 shows qualitative depth maps on five SCARED keyframes, confirming that the proposed pipeline recovers coherent tissue-surface geometry across the dataset.

765 4.7. Evaluation on the Stereo Laparoscopic Dataset

To evaluate generalization across diverse endoscopic hardware and surgical scenes, we further benchmark on a rectified stereo laparoscopic dataset comprising 22 sequences captured with five different endoscopes (focal lengths ranging from 383 to 766

Table 8: Aggregate evaluation on the Stereo Laparoscopic Dataset (22 sequences, 30 frames each, five endoscopes). Depth metrics (AbsRel, RMSE, δ_1) computed after median-scale alignment against GT depth maps. λ_3/λ_1 : avg 3D shape ratio (≥ 0.25 Normal 3D; $0.15\text{--}0.25$ Flat; < 0.15 Cone). **Bold**: best per metric.

Method	Normal 3D	Flat	Cone	λ_3/λ_1	AbsRel↓	δ_1 ↑	Time (s)
Mono	0	1	21	0.072	0.979	0.207	0.5
Stereo SGM	5	11	6	0.196	0.097	0.880	1.3
Ours	5	12	5	0.194	0.097	0.880	1.8

pixels, image resolutions from 360×288 to 720×576). Each sequence provides rec-
770 tified left/right image pairs with ground-truth depth maps; camera intrinsics and stereo
extrinsics (baseline ≈ 5.4 mm) are provided per sequence. We uniformly sample 30
stereo pairs per sequence and evaluate Ours, Stereo SGM, and Mono on depth accuracy
(AbsRel, RMSE) and 3D shape quality (λ_3/λ_1) against the provided GT depth.

Results. Table 8 summarizes the aggregate performance across all 22 sequences (per-
775 sequence breakdowns and cross-dataset metric plots are detailed in Supplementary
Material, Section G). As on SCARED, Ours and Stereo SGM achieve *identical* depth
accuracy (AbsRel = 0.097, δ_1 = 0.880), because the depth source is the same SGM
disparity; the role of the heightfield here is to attach a per-pixel confidence map rather
than to improve depth accuracy, and we therefore claim no depth-accuracy gain over
780 SGM. Both methods substantially outperform the monocular baseline (AbsRel = 0.979).
In terms of 3D shape quality the two are again comparable: Ours produces Normal 3D
geometry on 5 sequences versus 5 for SGM. What the experiment demonstrates is
hardware generalization: across five different endoscope configurations (focal lengths
383–766 px, resolutions 360×288 to 720×576) the heightfield maintains consistent
785 depth accuracy and confidence-map quality without re-tuning. Fig. 9 shows qualita-
tive depth maps on three representative sequences, confirming coherent tissue-surface
geometry.

4.8. Evaluation on Newly Captured Endoscopic Sequences

To test generalization beyond the bench-top chessboard and the public datasets, we
790 captured four new monocular endoscopic sequences with a different stereo-endoscope

Qualitative depth comparison on the Stereo Laparoscopic dataset (five sequences). The proposed SGM-heightfield pipeline recovers tissue-surface depth, whereas the monocular baseline yields structureless depth. (Turbo colormap, normalised per panel; predictions median-aligned to the GT for display.)




Figure 9: Qualitative depth comparison on the Stereo Laparoscopic dataset (three sequences). The proposed SGM-heightfield pipeline recovers coherent tissue-surface depth, whereas the monocular baseline yields structureless depth. (Turbo colormap, normalised per panel; predictions median-aligned to the GT for display.)

Cross-hardware evaluation on newly captured endoscopic sequences




Figure 10: Cross-hardware evaluation on the four newly captured sequences. (a) 3D shape ratio: Ours produces the best geometry on three of four sequences and is the only method reaching Normal 3D ($\lambda_3/\lambda_1=0.657$); on the longest sequence (V3) ORB-SfM is marginally better. (b) Metric scale fidelity (log scale): Ours stays within $\approx 2\times$ of the true scale on V1–V2, whereas ORB-SfM drifts 2–10 $\times$ and VGGT by 158–916 $\times$ (its bars exceed the axis). (c) ATE: VGGT is lowest but on 6–15 frames that overlap both the 60-frame subsample and the PnP reference; Ours is comparable to ORB-SfM on the full sequences.

rig (using the left channel only), at 1280 $\times$ 720 with calibrated intrinsics ($f_x=1210.7$, $f_y=1194.6$ px, distortion-corrected) and per-frame ground-truth poses from a 12 $\times$ 9 chessboard with 3.0 mm squares (PnP, reprojection RMSE <0.7 px). The four sequences (P1, V1, V2, V3; 496–839 frames, 30 fps) have intermittent board visibility (pose-valid
795 on 12–28% of frames; V1/V2/P1/V3 = 24%/28%/26%/12%) and tissue-like texture with specularities. Two smaller boards appear in the same recordings (9 $\times$ 6–5 mm on 3/5/19/0 frames; 9 $\times$ 8–8 mm on 0/1/1/1 frames) but are too sparse for a second ATE protocol, so all metric numbers use the 12 $\times$ 9–3 mm target (mean camera-to-board distance 177–216 mm; Supplementary Material, Section H.4). After removing
800 planar-PnP flip outliers (and, on P1, residual steps $\geq 60^\circ$), the mean relative rotation per valid-frame step is 1.9–4.2 $^\circ$ versus 0.3 $^\circ$ per frame on the chessboard benchmark; unfiltered means of 5–32 $^\circ$ are dominated by 180 $^\circ$ IPPE flips and are not used as motion statistics. We run the full classical pipeline (self-estimated poses) and evaluate against the cleaned PnP reference; ORB-SfM and VGGT are run under the same protocol
805 (VGGT on a 60-frame stride subsample due to GPU memory; 6–15 frames overlap the reference), and COLMAP is included as the classical SfM reference.

Results. Table 9 compares methods across all metrics. On geometric shape the proposed method is best on three of four sequences and is the only one reaching Normal 3D

Table 9: Cross-hardware evaluation on four newly captured endoscopic sequences (P1, V1, V2, V3) against independent PnP reference. Valid: n_{GT}/n_{frames} on 12×9–3 mm board. Dist.: mean camera-to-board distance (mm). λ_3/λ_1 : 3D shape ratio (≥ 0.25 Normal 3D; 0.15–0.25 Flat; < 0.15 Cone). ATE in mm; Scale is Sim3 factor (ideal 1.0). Corner depth: 11×8 inner corners after cloud alignment. COLMAP failed on all four sequences due to textureless video and is omitted. **Bold**: best per sequence per metric.

Seq	Valid	Dist.	Method	λ_3/λ_1	Shape	ATE↓	Scale	AbsRel	RMSE	δ_1
V1	24%	187	Ours	0.657	Normal 3D	27.5	0.47	0.065	15.4	0.978
			ORB-SfM	0.029	Cone	28.8	2.06	—	—	—
			VGGT	0.056	Cone	10.6	193	—	—	—
V2	28%	199	Ours	0.221	Flat	58.4	0.63	0.050	13.5	0.993
			ORB-SfM	0.011	Cone	68.1	9.86	—	—	—
			VGGT	0.063	Cone	48.9	916	—	—	—
P1	26%	216	Ours	0.119	Cone	55.2	0.35	0.070	19.8	0.991
			ORB-SfM	0.009	Cone	61.7	9.13	—	—	—
			VGGT	0.100	Cone	40.8	575	—	—	—
V3	12%	177	Ours	0.053	Cone	38.8	0.099	0.020	4.79	1.000
			ORB-SfM	0.245	Flat	25.1	4.32	—	—	—
			VGGT	0.012	Cone	8.2	158	—	—	—

(V1, $\lambda_3/\lambda_1=0.657$, 350k points). Metric scale follows the chessboard pattern: Ours
810 recovers scale 0.47 and 0.63 on V1–V2 (factors of $2.1\times$ and $1.6\times$), whereas ORB-SfM
drifts by $2\text{--}10\times$ and VGGT by $158\text{--}916\times$. COLMAP fails to reconstruct any of the four
sequences. Working distance is 177–216 mm on the trajectory board.

Failure modes.. Two sequences collapse to a Cone. V3 is the longest (839 frames) with
the sparsest board visibility (12% of frames) and a $10\times$ scale drift, consistent with weak
815 metric anchoring over long stretches. P1 is *not* a fast-rotation sequence once planar-
PnP flips are removed (mean 1.9° per valid-frame step for steps $<60^\circ$); the unfiltered
mean of 32° is dominated by 180° IPPE flips and is not a camera-motion rate. P1 is
instead the sequence whose PnP reference is noisiest (33 of 129 poses dropped). A
rotation-stratified breakdown (Supplementary Material, Section H) still shows a higher
820 position error in P1’s $>10^\circ$ bin (≈ 91 mm vs. ≈ 40 mm at $0\text{--}2^\circ$), but that bin has only
eight frames and may include residual flips; P1 errors are already large (≈ 40 mm) in
the slow-rotation bins. On V3 the same bins are nearly flat ($32\text{--}40$ mm), so the Cone
collapse is not explained by per-step rotation. Affine-aligned corner depth on V3 is low
(AbsRel 0.020, 1710 corners) after Sim3 alignment of the cloud; that local board-plane
825 fit does not restore the global $\lambda_3/\lambda_1=0.053$ shape. On V3, ORB-SfM’s incremental
chaining yields a marginally better shape (Flat, 0.245 vs. 0.053). These two failures
therefore support the stated limits on long, weakly anchored sequences; they do not by
themselves isolate a 32° rotation-collapse mechanism.

4.9. Ablation Study

830 Table 10 and Fig. 11 show the ablation study on key components of our pipeline.
Extended sensitivity results on the grid step and the search radius are provided in the
Supplementary Material, Section E.

ROI mask.. Removing the mask ($M_t \equiv 1$) lets background tissue gradients dominate
the heightfield, dropping λ_3/λ_1 from 0.42 to 0.18 ($\Delta=-0.24$, Flat), the largest single
835 degradation among all ablations.

Table 10: Ablation study on key pipeline components on the chessboard 50-frame set. The full-pipeline row matches Table 2. The three variant rows come from the same pipeline-of-record run; the intermediate clouds were not re-exported in this revision. $\Delta\lambda_3/\lambda_1$ denotes the absolute change in PCA shape ratio relative to the full pipeline.

Configuration	λ_3/λ_1	Δ	Shape	Points
Full pipeline	0.42	—	Normal 3D	11,735
Without ROI mask	0.18	−0.24	Flat	8,421
Without epipolar constraint	0.29	−0.13	Normal 3D	14,892
Without multi-view check	0.38	−0.04	Normal 3D	16,803




Figure 11: Ablation study visualization. (a) PCA shape ratio λ_3/λ_1 for each variant: removing the ROI mask drops the ratio below the Normal 3D threshold (0.25, green dashed), resulting in a Flat reconstruction. (b) Point cloud density: removing the multi-view check yields the most points but at lower shape quality; removing the mask yields the fewest. (c) Eigenvalue distribution (λ_2/λ_1 vs. λ_3/λ_1): the ROI mask primarily affects λ_3 (depth axis), while other components have more balanced effects. The figure shows all configurations tested during development; Table 10 reports the principal components. Delta values in (a) show change from Full pipeline.

Epipolar constraint.. Without the epipolar constraint, λ_3/λ_1 drops to 0.29 ($\Delta=-0.13$). Point count rises (11,735 $\rightarrow$ 14,892) but shape quality does not improve: the added points lie predominantly in high-gradient regions that already match well, while ambiguous matches on repetitive patterns go unchecked.

840 *Multi-view consistency check..* Without multi-view filtering, λ_3/λ_1 degrades to 0.38 ($\Delta=-0.04$) and point count increases (11,735 $\rightarrow$ 16,803) due to retained spurious matches from specular reflections.

4.10. Learned Stereo Matching Variants

To analyze learned stereo alternatives, we also evaluated five neural network variants
845 (V2–V4) trained on EndoNeRF (174 samples) and tested on zero-shot SCARED. The lightweight custom encoder (V2, 4.76M) attained in-domain AbsRel 0.0305 (matching training-free SGM), while over-parameterized models (V3, 26.5M) overfitted (AbsRel > 0.17). Pretrained ResNet backbones (V4) generalized better to zero-shot SCARED (AbsRel 2.435). Full architecture definitions, training curves, and Pareto frontier
850 analyses are detailed in Supplementary Material, Section E.

4.11. Cross-Dataset 3D Point-Cloud Geometry

To examine whether the well-formed 3D geometry observed on the chessboard benchmark generalises to surgical data, we reconstruct point clouds with the proposed pipeline and two references (classical Stereo SGM and the monocular gradient-inverse
855 baseline) on three additional stereo datasets (SCARED, EndoNeRF, and Stereo-Lap). Each point cloud is characterised by the PCA eigenvalue ratio λ_3/λ_1 , where a well-distributed 3D structure satisfies $\lambda_3/\lambda_1 \geq 0.25$ (Normal 3D); smaller values indicate Flat (0.15–0.25) or Cone (< 0.15) degeneracy. Table 11 reports the results (detailed 3D projection plots are in Supplementary Material, Section F).

860 The well-formed Normal 3D geometry is confined to the wide-baseline chessboard benchmark; on SCARED, Stereo SGM and the proposed variant both average just below the Normal 3D threshold (0.24 and 0.23 respectively across the eleven keyframes, Flat on average, Table 7); the confidence-based weighting slightly compresses the depth spread relative to plain SGM but does not preclude well-formed geometry on the

Table 11: PCA-based geometric shape of reconstructed point clouds across four datasets. λ_3/λ_1 is the ratio of the smallest to largest eigenvalue; Normal 3D (≥ 0.25), Flat (0.15–0.25), Cone (< 0.15). The chessboard benefits from a wide baseline and known planar geometry; the narrow-baseline surgical datasets (EndoNeRF, Stereo-Lap) limit depth resolution, so all methods degenerate there.

Dataset	Ours	SGM	Mono	Type
Chessboard	Norm.3D (0.42)	—	—	calib.
SCARED	Flat (0.23)	Flat (0.24)	Cone (0.11)	stereo
EndoNeRF	Flat (0.23)	Flat (0.23)	Flat (0.18)	narrow
Stereo-Lap	Flat (0.194)	Flat (0.196)	Cone (0.072)	stereo

Qualitative depth comparison on EndoNeRF. Three frames from each of the cutting and pulling sequences. The proposed calibrated end-to-end stereo variant recovers tissue surface depth consistent with the ground truth. Stereo SGM yields comparable but noisier depth, and the monocular baseline produces grossly inaccurate, structureless depth. (Color channels: normalized per panel; predictions median-aligned to the GT for display)




Figure 12: Qualitative depth comparison on EndoNeRF. Rows: three frames from the cutting sequence and three from the pulling sequence. Columns: RGB input, ground-truth depth, the proposed calibrated end-to-end stereo variant, Stereo SGM, and the monocular gradient-inverse baseline (Turbo colormap, normalised per panel; predictions median-aligned to the GT for display). The proposed method recovers tissue-surface depth consistent with the ground truth, Stereo SGM yields comparable but noisier depth, and the monocular baseline produces grossly inaccurate, structureless depth.

Table 12: Computational cost comparison. All methods were evaluated on the same hardware platform. [†]: includes global alignment optimization.

Method	Hardware	Time/50 fr	GPU Mem
ORB-SfM	CPU	8.2 s	—
COLMAP	CPU	42.5 s	—
VGGT	GPU	35.1 s	18.4 GB
Reloc3R	GPU	24.7 s	12.1 GB
DUS3R [†]	GPU	67.3 s	14.8 GB
MAS3R [†]	GPU	72.8 s	15.2 GB
Ours (Classic)	CPU	61.5 s	—
Ours-BA	GPU	15.8 s	4.2 GB

865 wider-baseline keyframes. On the narrow-baseline surgical datasets (EndoNeRF and Stereo-Lap), all three methods fall below the Normal 3D threshold, with mean shape ratios of 0.18–0.23 (Flat) for the stereo methods and 0.07–0.19 for the monocular baseline (Table 11): the small inter-camera baseline yields a disparity range of only a few pixels (EndoNeRF disparities span 41–43), limiting depth resolution and compressing
870 the reconstructed scene into a thin slab that no post-hoc processing can expand. Geometric well-formedness therefore requires sufficient baseline-induced parallax, and the chessboard result should not be extrapolated to arbitrary narrow-baseline surgical footage. Depth *accuracy* (Table 6) is preserved across these datasets because it depends on disparity estimation rather than absolute depth spread; only the 3D *shape* statistic
875 degrades under narrow-baseline geometry. Fig. 12 provides qualitative depth maps on the EndoNeRF cutting and pulling sequences, where the proposed calibrated variant recovers tissue-surface depth consistent with the ground truth, Stereo SGM yields comparable but noisier depth, and the monocular baseline produces structureless depth.

4.12. Computational Cost

880 Table 12 and Fig. 13 compare the computational cost. Our classical CPU-only implementation processes 50 frames in 61.5 s (~ 1.2 s per frame pair) with zero GPU




Figure 13: Computational cost comparison. (a) Runtime per 50 frames (logarithmic scale): blue bars indicate CPU-only methods; purple bars indicate GPU-accelerated methods. Per-frame processing time is annotated above each bar. (b) GPU memory consumption: all GPU methods require 12–18 GB; our method runs entirely on CPU. (c) Efficiency–quality trade-off: λ_3/λ_1 (geometric quality) vs. runtime. The dashed line marks the Normal 3D threshold (0.25); green shading highlights the ideal region (fast runtime with Normal 3D shape). Our method is the only approach in this region.

overhead, while the GPU-accelerated multi-view BA mode completes in 15.8 s using only 4.2 GB VRAM, substantially lower than VGGT (18.4 GB) and DUST3R (14.8 GB).

4.13. Comprehensive 17-Sequence Surgical Video Benchmark

To rigorously evaluate our framework across diverse endoscopic imaging dynamics, working distances, and tissue appearances, we construct an extensive benchmark spanning 17 continuous surgical video sequences (30 uniformly subsampled keyframes per sequence):

- **13 Real Surgical Sequences:** Captured with a calibrated laparoscopic camera observing surgical phantoms and calibration patterns under various standoff distances (89–260 mm), sweeping velocities, and viewing inclinations, categorized by protocol into test, validation, training, and extra-training splits.
- **4 Uncalibrated Simulation Cavities:** Rendered surgical cavities with realistic porcine soft tissue deformability, tissue folds, specular reflections, and dynamic surgical smoke, containing no physical calibration patterns.

We evaluate five representative approaches under identical protocols using the Absolute Trajectory Error (Sim3 ATE, mm) and reconstructed metric scale: (1) VGGT-1B [5], the 1.2B-parameter vision geometry transformer; (2) ORB-SfM [6], the classical

monocular tracking baseline; (3) Ours-BA, our multi-stride co-visibility Bundle Adjust-
 900 ment (Section 3.5) operating without physical landmarks; (4) Ours-SparseGauge, our
 multi-stride BA with metric scale anchored on a sparse keyframe subset (16 anchors, 14
 intermediate frames strictly held out; Section 3.6); (5) Ours-FullGauge, our complete
 framework utilizing metric gauge constraints across all visible frames.

Benchmark analysis... Table 13 and Fig. 14 document the benchmark results:

- 905 1. **Marker-free BA.** On all 17 sequences, Ours-BA has mean ATE 35.52 mm and
 median 29.88 mm. That median beats ORB-SfM (40.29 mm) and does not beat
 VGGT-1B (16.86 mm). On the four simulations the BA mean is 32.71 mm, VGGT-
 1B is 28.26 mm, and ORB-SfM is 35.75 mm. The single cavity `sim_s10097` is the
 exception versus VGGT-1B (24.86 vs. 37.25 mm); ORB-SfM is lower still on that
 910 clip (21.04 mm). An earlier single-clip run on `s10095` (37.72 vs. VGGT 38.36 and
 DUST3R 45.77 mm) is not the 17-sequence table and is not used below.
2. **Full gauge on the 13 real sequences.** FullGauge median ATE is 4.49 mm.
 The matching VGGT-1B and ORB-SfM medians on those same 13 sequences are
 11.88 mm and 41.98 mm. The 16.86 mm VGGT median quoted elsewhere is over all
 915 17 sequences, including the four simulations, where FullGauge is not defined. The
 4.49 mm figure therefore must not be compared with 16.86 mm. FullGauge uses the
 visible chessboard; VGGT-1B does not.
3. **Sparse Gauge Robustness:** In the practical sparse gauge setting (Ours-
 SparseGauge, with only 16 gauge frames and 14 held-out frames), the propagated
 920 metric scale bounds median ATE to 35.14 mm with near-unity scale factor ($1.035\times$ –
 $1.101\times$), confirming that metric scale is successfully disseminated across the co-
 visibility graph.
4. **Deep Pose Fusion Convergence:** As illustrated in Fig. 14(c), training our deep
 geometric pose fusion model across 11 training sessions over 120 epochs yields
 925 steady convergence, reaching a held-out validation ATE of 12.94 mm (direct position
 error 11.69 mm), closely matching VGGT-1B (12.76 mm) while training on less than
 0.1% of VGGT’s data volume.

Table 13: Comprehensive 17-sequence 3D reconstruction benchmark on real and simulated endoscopic video. All methods are evaluated on identical 30-frame keyframe sequences. ATE: Absolute Trajectory Error after Sim3 alignment (mm, ↓). Ours-BA: Multi-stride bundle adjustment without landmark tracking. Ours-SparseGauge: Metric scale anchored on sparse keyframes (stride 2; 16 anchors, 14 held-out frames). Ours-FullGauge: Complete framework with full metric gauge injection. Simulation sequences contain no physical calibration patterns (gauge modes marked as “—”). The 13-real median uses only the real sequences, where FullGauge is defined. The 17-seq median uses every sequence; FullGauge is blank there. **Bold**: lowest ATE in that row.

Sequence	Kind	Split	VGGT-1B	ORB-SfM	Ours-BA	Ours-SparseG	Ours-FullG
traj_023422	real	test	16.86	68.59	65.87	63.97	4.55
traj_144019	real	val	12.76	65.99	43.36	49.58	4.66
traj_023241	real	val	18.95	71.39	32.37	47.87	15.90
traj_015507	real	train	43.54	56.66	74.48	71.54	80.25
traj_015629	real	train	28.75	35.53	21.04	35.14	26.38
traj_015918	real	train	5.70	29.52	24.15	19.23	1.31
traj_020132	real	train	21.83	35.48	29.88	30.94	2.35
traj_020407	real	train	11.88	41.98	27.87	24.95	1.85
traj_143156	real	train	9.35	58.75	58.66	49.38	3.12
rigid_123958	real	extra	4.14	21.58	17.84	21.04	4.49
rigid_124243	real	extra	7.30	32.47	18.48	20.38	3.44
rigid_140510	real	extra	7.53	45.00	36.33	40.52	8.51
rigid_142902	real	extra	5.99	37.33	22.67	34.20	2.05
sim_s10095	sim	sim	36.97	42.13	39.85	—	—
sim_s10096	sim	sim	18.07	40.29	41.60	—	—
sim_s10097	sim	sim	37.25	21.04	24.86	—	—
sim_s10098	sim	sim	20.75	39.55	24.53	—	—
13-real median	—	—	11.88	41.98	29.88	35.14	4.49
17-seq median	—	—	16.86	40.29	29.88	—	—




Figure 14: Comprehensive 17-sequence surgical video benchmark and deep pose training dynamics. (a) Sim3 ATE across representative real surgical test, validation, and training recordings, as well as smoke-filled simulation cavities. (b) Medians from Table 13. FullGauge 4.49 mm is the median on the 13 real sequences; the VGGT-1B median on those same 13 sequences is 11.88 mm. The 16.86 mm VGGT median is over all 17 sequences and is not the matched comparison. Marker-free BA median on all 17 is 29.88 mm. (c) 120-epoch training convergence of the deep geometric pose fusion network (ResNet-50 backbone with spatial attention pooling), demonstrating smooth loss reduction and convergence to 12.94 mm validation ATE on the held-out validation sequence. (d) Reconstructed Sim3 scale distribution (log scale): Ours-FullGauge tightly tracks the ideal $1.0\times$ physical scale ($0.998\times-1.019\times$), whereas unconstrained baselines suffer from $148\times-325\times$ scale distortion.

4.14. Protocol-Aware Summary

The evaluations above answer different questions and must not be collapsed into a single ranking. Table 14 collects the headline numbers under the protocol that produced them; each row is an independent comparison. We explicitly distinguish two tracking questions. (1) On the short chessboard bench (Table 2), VGGT has the lower ATE (3.02 vs. 3.89 mm) and the lower corner AbsRel (0.009 vs. 0.247), while Ours keeps Normal 3D shape ($\lambda_3/\lambda_1 = 0.42$) and a Sim3 scale of $1.11\times$ (VGGT $91\times$). (2) On the 17-sequence table, FullGauge and VGGT answer different questions. FullGauge median 4.49 mm is over the 13 real sequences with a visible board; the VGGT median on those same 13 sequences is 11.88 mm, not the 17-sequence median 16.86 mm. Marker-free BA does not beat VGGT on the 17-sequence median (29.88 vs. 16.86 mm). Fig. 15 and Table 14 put each comparison on one row. A green cell is a win for Ours on that row only. Second, on public EndoNeRF, the calibrated stereo variant attains the lowest median-aligned AbsRel (0.0285), ahead of FoundationStereo (0.0321) and Stereo SGM (0.0323). On SCARED and the stereo-laparoscopic set the classical pipeline matches SGM in depth by construction and adds a confidence map; we do not claim a depth-accuracy gain there. On the newly captured sequences the same scale pattern recurs on V1 ($0.47\times$ vs. VGGT $193\times$) and V2 ($0.63\times$ vs. VGGT $916\times$). Native square spacing, with no depth affine, is a separate question: our native median edge is 5.49 mm against the 5.0 mm square, and the mean absolute error is 10.9 mm (Table 3). With a sparse baseline gauge at triangulation, the hold-out median MAE falls to 1.45 mm, below COLMAP (3.41 mm), VGGT (4.65 mm), and ORB-SfM (4.85 mm).

5. Discussion

Why feed-forward networks fail on endoscopic data.. The domain gap between Internet-scale natural-image training sets and endoscopic imaging causes two specific failures: (1) specular reflections and smooth tissue provide few reliable features, breaking correspondence estimation; (2) repetitive patterns (e.g., calibration grids, tissue textures) confuse learned matching priors, leading to incorrect depth ordering and collapsed 3D structure. Our method mitigates both through gradient-based heightfield construction

Table 14: Protocol-matched headline results. Each row uses one metric and one set of sequences. Bold marks the better entry on that row. FullGauge is defined only when the chessboard is visible, so its median is over 13 real sequences, not over all 17.

Setting	Question	Ours	Strongest baseline		Outcome
13 real, board visible	Median ATE	4.49 mm (FullG)	VGGT 41.98	11.88; ORB	Better, with the board
17 seq., no board	Median ATE	29.88 mm (BA)	VGGT 40.29	16.86 ; ORB	Better than ORB only
4 simulations	Mean ATE	32.71 mm (BA)	VGGT 35.75	28.26 ; ORB	Better than ORB only
Sim s10097	ATE, no gauge	24.86 mm (BA)	VGGT 21.04	37.25; ORB	Better than VGGT only
Chessboard	Shape λ_3/λ_1	0.42 (Normal)	SGM 0.31; else ≤ 0.22		Best; SGM also Normal
Chessboard	Sim3 scale vs PnP	1.11 $\times$	ORB 3.7 $\times$; VGGT 91 $\times$		Best metric scale
Chessboard	ATE vs PnP (mm)	3.89	VGGT 3.02		Second
Chessboard	Corner AbsRel	0.247	VGGT 0.009		Not best
Chessboard	Square spacing	med. 5.49, MAE 10.9	ORB med. 6.26, MAE 5.89		Near 5 mm; large error
Chessboard	Gauge square MAE	1.45	COLMAP 3.41; VGGT 4.65		Lowest (1.45 mm)
EndoNeRF-100	AbsRel (median)	0.0285	FS 0.0321; SGM 0.0323		Best among compared
SCARED / Lap	Depth vs SGM	identical	SGM (same disparity)		Tie by construction
New V1	Shape / scale	0.657 / 0.47 $\times$	VGGT 0.056 / 193 $\times$		Best shape+scale
New V2	Shape / scale	0.221 (Flat) / 0.63 $\times$	VGGT 916 $\times$		Best scale; Flat
New P1	Shape / scale	0.119 (Cone) / 0.35 $\times$	VGGT 575 $\times$		Collapse
New V3	Shape / scale	0.053 (Cone) / 0.099 $\times$	ORB Flat 0.245 / 4.3 $\times$		Collapse




Figure 15: Protocol-matched comparison. (a) Trajectory ATE under three groupings taken from Table 13. “13 real with board” compares FullGauge with VGGT-1B and ORB-SfM on the same 13 sequences. “17 seq. no board” is marker-free BA. “4 sim mean” is the mean over the four simulations. (b) Other tasks. Green means Ours is better on that metric; red means the named baseline is better. The scale row compares $1.11\times$ with VGGT’s $91\times$.

(uniform specular responses are suppressed by percentile truncation) and the ROI mask (constraining reconstruction to geometrically reliable regions); the same heightfield geometry is also the substrate on which inter-frame poses are estimated, so reliable localization and reliable correspondences share a common 3D signal that 2D-only or learned-natural-image methods lack. Across the chessboard, public stereo sets, and the newly captured sequences, no feed-forward method achieves $\lambda_3/\lambda_1 \geq 0.25$ on any sequence tested.

Generalization across imaging hardware.. The stereo laparoscopic evaluation (Section 4.7) demonstrates that our method generalizes across five different endoscope configurations (focal lengths 383–766 px, resolutions 360×288 to 720×576) without re-tuning, achieving depth accuracy identical to stereo SGM while providing an additional heightfield confidence map. This hardware-agnostic behavior stems from the gradient-based design: percentile truncation adapts automatically to the gradient distribution of each frame, making the heightfield independent of absolute image intensity or resolution.

Implications for imaging, visualization, and image-guided intervention.. From an imaging-and-graphics standpoint, the useful output is not a disparity map alone but

a surface that can be inspected, overlaid, and measured. The point cloud—whose poses
 975 are estimated from the heightfield rather than supplied externally—can be rendered and
 overlaid. A lesion-diameter or standoff measurement is not supported by the present
 square-size error: on the 5.0 mm board the native mean absolute edge error is 10.9 mm
 (Section 4.3). Overlay and measurement both fail when the cloud collapses to a cone
 or plane. Relative to electromagnetic tracking (sensor coils, metal interference) and
 980 intra-operative CT or ultrasound (added hardware or dose), a vision-only pipeline that
 reuses the existing endoscope feed adds no device to the sterile field. The present vali-
 dation is computational rather than clinical: it uses ex-vivo porcine tissue (SCARED),
 de-identified surgical recordings (EndoNeRF), and a bench-top chessboard phantom,
 with no live patients or prospective surgical use. The chessboard Normal 3D result re-
 985 quires a sufficiently strong heightfield signal; the ablation in Section 4.3 shows that a 2D
 pose front-end collapses the cloud, so immediate deployment is limited to settings with
 adequate texture and viewing geometry (calibrated bench-top and stereo-laparoscopic
 sequences) rather than arbitrary marker-free footage. Marker-free masks via instru-
 ment and tissue segmentation, and a prospective phantom or cadaver study against
 990 electromagnetic tracking, are the natural next steps toward clinical translation.

Limitations.. First, the default ROI is a landmark hull or a contour fallback; on the new
 sequences the board is visible on only 12–28% of frames, so a semantic tissue mask
 (Section 3.3) is the intended replacement rather than a second reconstruction primitive.
 Second, the 2.5-D heightfield cannot represent occlusions or multiple depth layers, and
 995 degrades when scene motion is predominantly lateral (Section 4.4). Instruments must
 be masked out, not lifted as a second surface. Third, pairwise heightfield-ICP accu-
 mulates drift on long, weakly anchored sequences (V3: scale 0.099 $\times$, $\lambda_3/\lambda_1=0.053$,
 ATE 38.8 mm). The 7-D matches were not archived, so that residual was not rerun
 on the published cloud. A pose-only substitute re-estimated every fourth relative pose
 1000 by Kabsch on high-gradient heightfield vertices and reintegrated the chain. On the 27
 frames of that stride that meet the PnP reference, Sim3 scale moved from 0.090 to
 0.081 and ATE from 35.7 mm to 33.6 mm. Scale did not approach 1. The published
 full-sequence figures are unchanged. Fourth, reconstruction remains confined to M_t ; a

soft semantic mask would enlarge coverage without changing the heightfield equations.

1005 Fifth, orientation in the heightfield embedding is only approximate. Translation and scale are reliable on the chessboard (ATE 3.89 mm, scale 1.11 $\times$) because that motion is translation-dominant; rotation-dominant footage needs the same BA stage, not a different heightfield. A detailed analysis of failure cases is provided in the Supplementary Material, Section F.3.

1010 6. Conclusion

We have presented a unified geometric reconstruction framework integrating pseudo-3D heightfield surface representation, normal-guided branch disambiguation, multi-stride bundle adjustment, and metric scale gauge injection for endoscopic video. The framework offers dual operating modes: a lightweight, CPU-only classical variant
1015 for resource-constrained clinical systems, and a GPU-accelerated multi-stride BA variant for high-precision global trajectory optimization. The physical heightfield normal selects one of the two homography branches (Section 3.5). On the 13 real sequences with a visible chessboard, full-gauge median ATE is 4.49 mm, against 11.88 mm for VGGT-1B and 41.98 mm for ORB-SfM. Marker-free BA does not beat VGGT-1B on
1020 the 17-sequence median (29.88 vs. 16.86 mm). On `sim_s10097`, BA is 24.86 mm and VGGT-1B is 37.25 mm; ORB-SfM is 21.04 mm on that clip. On public EndoN-erF, the calibrated stereo variant reaches AbsRel 0.0285 (FoundationStereo 0.0321). The heightfield’s separate claim is the un-collapsed surface: $\lambda_3/\lambda_1 = 0.42$ on the chessboard, where the feed-forward clouds are Flat or Cone.

1025 Acknowledgements

The authors thank the developers of the SCARED [43] and EndoNeRF [38] datasets for making their data publicly available, which made the cross-dataset evaluation possible.

Declarations

Funding. This work was supported by the National Natural Science Foundation of China (Grant No. 6240074117); the Natural Science Project of the Science and Technology Department of Henan Province (Grant Nos. 232102211001, 232102210042, 232102210005, 222102310058); the Open Project Program of the State Key Laboratory of Virtual Reality Technology and Systems (Grant No. VRLAB2023C03); the High-level Talents Fund of Henan University of Technology (Grant Nos. 2022BS076, 2022BS075); and the Cultivation Project of Tuoxin Team in Henan University of Technology (Grant No. 2024TXTD18).

Conflict of interest. The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Ethics statement. This study uses public datasets and in-house bench-top phantom recordings. The SCARED dataset [43] contains ex-vivo porcine cadaver tissue and does not involve human subjects. The EndoNeRF dataset [38] consists of de-identified surgical recordings previously released by the original authors under their own institutional ethics approval; the present study is a secondary computational analysis of these public recordings and did not require additional ethical clearance. The in-house endoscopic chessboard, stereo-laparoscopic, and newly captured monocular sequences are bench-top phantom recordings and contain no patient-identifiable information.

Authors' contributions. **Ranyang Li:** Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Visualization, Funding acquisition. **Wufeng Liu:** Methodology, Validation, Writing – review & editing. **Xiaodong Guo:** Software, Validation, Data curation. **Nan Wei:** Resources, Supervision. **Chao Fan:** Software, Investigation. All authors read and approved the final manuscript.

Data availability. Public datasets are cited at their original sources: SCARED [43] and EndoNeRF [38]. Evaluation scripts, the pipeline of record, result JSONs, and the figures

used in this manuscript are at <https://github.com/liranyang123456-commits/re3d-heightfield>. In-house sequences and trained stereo-variant weights will be deposited in the same repository; until the full data release they remain available from the corresponding author upon reasonable request. The random seed used for PCA sub-sampling is 42.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process.. During the preparation of this work the authors used generative-AI tools solely for language polishing of author-written text. After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

References

- [1] D. Stoyanov, G. P. Mylonas, F. Deligianni, A. Darzi, G.-Z. Yang, Soft-tissue motion tracking and structure estimation for robotic assisted MIS procedures, in: MICCAI, 2005, pp. 1–8.
- [2] P. Mountney, G.-Z. Yang, Simultaneous estimation of 3D structure and motion for endoscopic navigation, IEEE Trans. Med. Imag. 29 (7) (2010) 1313–1324.
- [3] S. Wang, V. Leroy, Y. Cabon, B. Chidlovskii, J. Revaud, DUS3R: Geometric 3D vision made easy, in: CVPR, 2024, pp. 20697–20709. doi:10.1109/CVPR52733.2024.01956.
- [4] V. Leroy, Y. Cabon, J. Revaud, Grounding image matching in 3D with MAST3R, in: ECCV, 2024, pp. 71–91.
- [5] J. Wang, M. Chen, N. Karaev, A. Vedaldi, C. Rupprecht, D. Novotny, VGGT: Visual geometry grounded transformer, in: CVPR, 2025, pp. 5294–5306. doi:10.1109/CVPR52734.2025.00499.
- [6] R. Mur-Artal, J. M. M. Montiel, J. D. Tardós, ORB-SLAM: A versatile and accurate monocular SLAM system, IEEE Trans. Robot. 31 (5) (2015) 1147–1163. doi:10.1109/TR0.2015.2463671.

- [7] J. L. Schönberger, J.-M. Frahm, Structure-from-motion revisited, in: CVPR, 2016, pp. 4104–4113. doi:10.1109/CVPR.2016.445.
- [8] R. A. Newcombe, et al., DTAM: Dense tracking and mapping in real-time, in: ICCV, 2011, pp. 2320–2327.
- [9] T. Whelan, R. F. Salas-Moreno, B. Glocker, A. J. Davison, S. Leutenegger, ElasticFusion: Real-time dense SLAM and light source estimation, in: RSS, 2015. doi:10.15607/RSS.2015.XI.001.
- [10] J. Reizenstein, et al., Common objects in 3D: Large-scale learning and evaluation, in: ICCV, 2021, pp. 10903–10912.
- [11] B. K. P. Horn, Obtaining shape from shading information, in: B. K. P. Horn, M. J. Brooks (Eds.), Shape from Shading, MIT Press, 1989, pp. 123–171.
- [12] D. Stoyanov, M. V. Scarzanella, P. Pratt, G.-Z. Yang, Real-time stereo reconstruction in robotically assisted minimally invasive surgery, in: MICCAI, 2010, pp. 275–282. doi:10.1007/978-3-642-15705-9_34.
- [13] D. Eigen, et al., Depth map prediction from a single image using a multi-scale deep network, in: NeurIPS, 2014, pp. 2366–2374.
- [14] R. Ranftl, et al., Vision transformers for dense prediction, in: ICCV, 2021, pp. 12179–12188.
- [15] K. B. Ozyoruk, G. I. Gokceler, T. L. Bobrow, G. Coskun, K. Incetan, Y. Almalioglu, F. Mahmood, E. Curto, L. Perdigoto, M. Oliveira, et al., EndoSLAM dataset and an unsupervised monocular visual odometry and depth estimation approach for endoscopic videos, Medical Image Analysis 71 (2021) 102058. doi:10.1016/j.media.2021.102058.
- [16] H. Hirschmüller, Stereo processing by semi-global matching and mutual information, IEEE Trans. Pattern Anal. Mach. Intell. 30 (2) (2008) 328–341. doi:10.1109/TPAMI.2007.1166.

- 1110 [17] F. Mourgues, E. Coste-Maniere, Flexible calibration of actuated stereoscopic
endoscope for overlay in robot assisted surgery, in: MICCAI, 2002, pp. 25–34.
- [18] E. Rublee, et al., ORB: An efficient alternative to SIFT or SURF, in: ICCV, 2011,
pp. 2564–2571.
- [19] D. G. Lowe, Distinctive image features from scale-invariant keypoints, *Int. J.*
1115 *Comput. Vis.* 60 (2) (2004) 91–110.
- [20] J. Engel, et al., LSD-SLAM: Large-scale direct monocular SLAM, in: ECCV,
2014, pp. 834–849.
- [21] R. A. Newcombe, et al., KinectFusion: Real-time dense surface mapping and
tracking, in: ISMAR, 2011, pp. 127–136.
- 1120 [22] A. Dai, et al., BundleFusion: Real-time globally consistent 3D reconstruction, in:
CVPR, 2017, pp. 5765–5773.
- [23] Y. Cabon, L. Stofl, L. Antsfeld, G. Csurka, B. Chidlovskii, J. Revaud, V. Leroy,
MUS3R: Multi-view network for stereo 3D reconstruction, in: CVPR, 2025, pp.
1050–1060.
- 1125 [24] S. Dong, S. Wang, S. Liu, L. Cai, Q. Fan, J. Kannala, Y. Yang, Reloc3R: Large-
scale training of relative camera pose regression for generalizable, fast, and accu-
rate visual localization, in: CVPR, 2025, pp. 16739–16752.
- [25] Z. Zhang, A flexible new technique for camera calibration, *IEEE Trans. Pattern*
Anal. Mach. Intell. 22 (11) (2000) 1330–1334. doi:10.1109/34.888718.
- 1130 [26] P. F. Sturm, S. J. Maybank, On plane-based camera calibration: A general algo-
rithm, *IEEE Trans. Pattern Anal. Mach. Intell.* 23 (9) (2001) 965–980.
- [27] Z. Zhao, S. Voros, Y. Weng, F. Chang, R. Li, Tracking-by-detection of surgical
instruments in minimally invasive surgery via the convolutional neural network
deep learning-based method, *Comput. Assist. Surg.* 22 (sup1) (2017) 26–35.
1135 doi:10.1080/24699322.2017.1378777.

- [28] G. Bradski, The OpenCV library, Dr. Dobb's J. Software Tools (2000).
- [29] E. Malis, M. Vargas, Deeper understanding of the homography decomposition for vision-based control, Inria Research Report 6303 (2007).
- [30] J. Sun, Z. Shen, Y. Wang, H. Bao, X. Zhou, LoFTR: Detector-free local feature
1140 matching with transformers, in: CVPR, 2021, pp. 8922–8931.
- [31] S. Woo, J. Park, J.-Y. Lee, I. S. Kweon, CBAM: Convolutional block attention module, in: ECCV, 2018, pp. 3–19.
- [32] X. Guo, K. Yang, W. Yang, X. Wang, H. Li, Group-wise correlation stereo network, in: CVPR, 2019, pp. 3273–3282.
- [33] J. R. Chang, Y. S. Chen, Pyramid stereo matching network, in: CVPR, 2018, pp.
1145 5410–5418.
- [34] C. Wang, D. Xu, Y. Zhu, R. Martín-Martín, C. Lu, L. Fei-Fei, S. Savarese, DenseFusion: 6d object pose estimation by iterative dense fusion, in: CVPR, 2019, pp. 3343–3352.
- [35] J. Sturm, et al., A benchmark for the evaluation of RGB-D SLAM systems, in:
1150 IROS, 2012, pp. 573–580.
- [36] D. Nistér, An efficient solution to the five-point relative pose problem, IEEE Trans. Pattern Anal. Mach. Intell. 26 (6) (2004) 756–770.
- [37] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, et al., DINOv2: Learning robust visual features without supervision,
1155 Trans. Mach. Learn. Res. (2024).
- [38] Y. Wang, Y. Long, S. H. Fan, Q. Dou, Neural rendering for stereo 3D reconstruction of deformable tissues in robotic surgery, in: MICCAI, 2022, pp. 431–441. doi : 10.1007/978-3-031-16449-1_41.
- [39] B. Wen, M. Trepte, J. Aribido, J. Kautz, O. Gallo, S. Birchfield, FoundationStereo: Zero-shot stereo matching, in: CVPR, 2025, pp. 5249–5260. doi : 10.1109/CVPR52734.2025.00495.

- 1165 [40] G. Xu, X. Wang, X. Ding, X. Yang, Iterative geometry encoding volume for stereo matching, in: CVPR, 2023, pp. 21919–21928. doi:10.1109/CVPR52729.2023.02099.
- [41] D. Scharstein, R. Szeliski, A taxonomy and evaluation of dense two-frame stereo correspondence algorithms, *Int. J. Comput. Vis.* 47 (1–3) (2002) 7–42.
- [42] L. Yang, B. Kang, Z. Huang, X. Xu, J. Feng, H. Zhao, Depth anything V2, in: NeurIPS, 2024.
- 1170 [43] M. Allan, J. McLeod, C. Wang, J. C. Rosenthal, Z. Hu, N. Gard, P. Eisert, K. X. Fu, T. McLeod, X. Luo, et al., Stereo correspondence and reconstruction of endoscopic data challenge, arXiv:2101.01133 (2021). doi:10.48550/arXiv.2101.01133.